



Systemic Risk Assessment

August 2025
Assessment Report

Table of Contents

1. Introduction	4
2. Executive Summary	5
3. Background	7
4. Products at a Glance.....	10
4.1 Overall Product Design	10
4.2 Existing Products	11
4.3 New since Last Year	32
5. Methodology	39
5.1 Year 3 Design Enhancements	39
5.2 Risk Assessment Process	40
5.3 Risk Assessment Tools.....	43
6. Summary of Results.....	44
6.1 Key Findings	44
Consumer Protection and Fraud.....	45
Human Dignity	49
Mental and Physical Wellbeing.....	52
Discrimination and Hate.....	55
Freedom of Expression and Information.....	59
Protection of Personal Data.....	61
Civic Discourse and Electoral Processes	64
Public Health	68
Illegal Content and Activities.....	69
Private and Family Life	72
Public Security	75
Rights and Protection of Minors.....	77
6.2 Scoring Rationale	81
7. Mitigations.....	85
7.1 Best Practice Highlights	85
7.2 Monitoring the Effectiveness of Mitigations.....	107
7.3 Progress Made on Enhanced Mitigations	108
7.4 Planned Enhanced Mitigation Efforts.....	109
8. Conclusion	110
9. Appendices	111

How to Read this Report

- For a description of LinkedIn products and services, see: [Products at a Glance](#)
- To review this year's risk assessment results, see: [Summary of Results](#)
- To read about changes since the last report, see: [New Since Last Year](#), [Year 3 Design Enhancements](#), [Year 3 Changes in Risk Scores](#), [Progress Made on Enhanced Mitigations](#), and [Planned Enhanced Mitigation Efforts](#)
- To understand how LinkedIn measures and monitors the effectiveness of its mitigation measures, see: [Residual Risk](#) and [Monitoring the Effectiveness of Mitigations](#)
- For details on the methodology LinkedIn used to conduct this risk assessment, see: [Methodology](#) and [Methodology Expanded](#)
- For an overview of the full risk mitigation measures LinkedIn has implemented, see: [Best Practice Highlights](#)

Jump to Findings by Risk Area

[1 Civic Discourse and Electoral Processes](#)

[2 Consumer Protection and Fraud](#)

[3 Discrimination and Hate](#)

[4 Freedom of Expression and Information](#)

[5 Human Dignity](#)

[6 Illegal Content and Activities](#)

[7 Mental and Physical Wellbeing](#)

[8 Private and Family Life](#)

[9 Protection of Personal Data](#)

[10 Public Health](#)

[11 Public Security](#)

[12 Rights and Protection of Minors](#)

1. Introduction

LinkedIn Ireland Unlimited Company—the provider of LinkedIn’s services in the European Union (EU)—has been designated by the European Commission as a Very Large Online Platform (VLOP) and is therefore subject to the Digital Services Act (DSA) Article 34 requirement to conduct a Systemic Risk Assessment. In accordance with DSA Article 42(4)(a), this Systemic Risk Assessment Report (the “Report”) sets out the following: the results of LinkedIn’s Article 34 risk assessment, the specific mitigation measures LinkedIn has put in place pursuant to Article 35(1), and information about the consultations conducted by LinkedIn in support of the risk assessment and design of risk mitigation measures.

The Report reflects the Year Three assessment, through August 2025, of systemic risks indicated by DSA Article 34 stemming from the design or functioning of LinkedIn or from the use of LinkedIn services in the EU. The assessment considered the specific risk categories and influencing factors outlined in Article 34 with a focus on their impact in the EU; however, LinkedIn also evaluated these risks and influencing factors at a global level. Accordingly, the findings and mitigations discussed in this Report are largely applicable and relevant to the other regions and countries in which LinkedIn offers its services.

The Report describes assessed risk related to the LinkedIn platform and services holistically. The Report also addresses specific risks or mitigations as they may apply uniquely to specific LinkedIn products. The Report includes an executive summary, background on LinkedIn as it relates to the Systemic Risk Assessment, a brief overview of LinkedIn products to inform their risk profiles, LinkedIn’s approach to and the methodology used for the Systemic Risk Assessment, a summary of the Systemic Risk Assessment results, and an overview of LinkedIn’s risk mitigation efforts, both in place and planned.

2.Executive Summary

LinkedIn is committed to keeping its platform and services safe, trusted, and professional and to providing transparency to its members, the public, and to regulators. LinkedIn's vision is to create economic opportunity for every member of the global workforce. Its mission is to connect the world's professionals to make them more productive and successful. Members come to LinkedIn to find a job, stay informed, connect with other professionals, and learn new skills. As a real-identity online networking service for professionals to connect and interact with other professionals, LinkedIn has a unique risk profile when compared with many social media platforms.

The systemic risks with the highest levels of potential residual risk following the application of risk mitigation measures on LinkedIn's platform are (1) Consumer Protection and Fraud and (2) Human Dignity. While LinkedIn has robust policies, detection capabilities, and enforcement processes in place to mitigate risk in these areas, due to the evolving and pervasive nature of the related threats, LinkedIn recognizes that continued investment will be required to continue to decrease these risks. Some examples include:

- LinkedIn prioritizes investment in strategies to mitigate and manage Consumer Protection and Fraud risks and continuously uses feedback to improve deficiencies. For example, LinkedIn's trust team identified an increase in fake account creation in December 2024, and after investigation, it was attributed to a session hijacking attack. To combat this risk, LinkedIn deployed a new mitigation to enhance verification requirements for free user trial sign-ups. Following the implementation of this safeguard, LinkedIn monitored a decrease in fake accounts created. Additionally, LinkedIn engaged with the EU Directorate-General for Communications Networks to obtain views of independent experts and discussed feedback on mitigating fraud and scams on the platform.
- LinkedIn expressly forbids sexual content, nudity and human trafficking as outlined in the User Agreement and the Professional Community Policies. If any such content appears on the platform, LinkedIn uses a three-layered approach to moderate it, leveraging both machine learning and human moderators. LinkedIn's cross-functional Trust Product team works together not only to detect and remove abusive content from LinkedIn, but also to educate LinkedIn members that unwanted sexual advances and other sexual harassment do not belong on LinkedIn. Additionally, LinkedIn scans job listings and advertisements for potential policy violations related to human trafficking and allows user reporting of content including jobs and ads.

While other risk areas examined maintain a lower level of assessed residual risk, LinkedIn continues to invest in numerous Trust domains to proactively enhance the safety, security, privacy, and quality of the LinkedIn user experience.

New and sustained LinkedIn investments include:

- The continued development and refinement of advanced models for detecting and preventing the creation of fake or inauthentic accounts, fraudulent activities, or sharing of harmful content, including Child Sexual Abuse Material (CSAM), hate speech, and harassment;
- Investment in human reviewer subject matter expertise, high quality training for Trust team members, and a robust quality assurance program for manual content moderation decisions;
- Sophisticated internal tooling for content moderation to drive accountability, quality assurance, cross-functional collaboration, and reporting; and
- Enhanced Professional Community Policies to explicitly prohibit synthetic and manipulated media created or manipulated to distort real life events to prevent spread of mis-and-disinformation.

3. Background

LinkedIn's vision is to create economic opportunity for every member of the global workforce. Its mission is to connect the world's professionals to make them more productive and successful. LinkedIn is a networking tool that enables members to establish their professional identities online, connect with other professionals, and build relationships for the purpose of collaborating, learning, and staying informed about industry information and trends. As such, the design and function of the platform are central to its overall risk profile, which they shape in a few key ways:

- LinkedIn is a real-identity platform, where members must use their real or preferred *professional* names, and the content they post is visible, for example, to their colleagues, employers, potential future employers, and business partners. Given this audience, members tend to limit their activity to professional areas of interest and expect the content they see to be professional in nature.
- LinkedIn operates under standards of professionalism, which are reflected both in content policies and enforcement, as well as in content prioritization and amplification. LinkedIn's policies bolster a safe, trusted, and professional platform, and LinkedIn strictly enforces them. LinkedIn strives to broadly distribute high-quality content that advances professional conversations on the platform.
- LinkedIn services are tailored toward professionals and businesses. During the account creation process, new members listed as attending a secondary education institution must attest to being over the minimum age (16 years) to open an account.
- LinkedIn's Digital Safety function ensures a "safety by design" approach throughout the product development lifecycle by partnering with the relevant product and engineering organizations to conduct continuous assessments of risks and address threats prior to product launch.

These factors significantly reduce the likelihood of certain systemic risks stemming from LinkedIn's platform and services. Nonetheless, LinkedIn proactively invests in establishing robust policies, advanced detection capabilities, and effective enforcement processes in line with the Digital Trust and Safety Partnership's (DTSP) Best Practices Framework. Additionally, as bad actors may resort to creating fake personas to attempt abusive or malicious activities on the platform, LinkedIn invests significant resources in developing techniques to prevent the creation of fake Profiles and provides identity verification and other authenticity signals to members.

One of LinkedIn's core values is "**members first.**" This means that every day, when making decisions large and small, LinkedIn asks "Is this the right thing to do for our members?" This value drives the company and pushes it to remain worthy of its members' trust worldwide.

As LinkedIn's impact on the world of work has grown, the company recognizes that its responsibility has grown as well. LinkedIn understands that in the critical areas of privacy, security, safety, and fairness, it must continue to deliver for members and the world beyond. So, in addition to its "members first" foundational value, LinkedIn operates under a set of "Guiding Principles for Trust" to make the best possible decisions to protect LinkedIn members and maintain their trust:

- LinkedIn provides its members with **clarity, consistency, and control** over their data. LinkedIn tells members what it will do with their data. LinkedIn then does what it says it will do with this data. And LinkedIn strives to provide its members with simple, easy to understand settings and controls so that they are in control of their data.
- LinkedIn is focused on keeping its platform **safe, trusted, and professional**. It uses systems, technology, and reports from its members to detect and quickly remove content that violates LinkedIn's Professional Community Policies. Members also rightfully expect content they encounter on LinkedIn to be legitimate. Removing fake accounts, jobs, and Pages is paramount to this objective.
- LinkedIn believes that members with equal talent should have **equal access** to opportunity. To achieve this goal, LinkedIn is committed to building products with no unfair bias.
- LinkedIn is a global platform with an obligation to **respect the laws** that apply to it. It also contributes to the dialogue that shapes these laws so that LinkedIn can fulfill its vision of creating economic opportunity for every member of the global workforce.

At LinkedIn, trust is the cornerstone for building meaningful connections and professional relationships. Members rely on LinkedIn to create an environment where they can safely learn and grow in their careers. To execute on this responsibility, LinkedIn takes a multidimensional approach with three layers of protection within its ecosystem to help it filter out content and behavior that violates its policies such as such as harassment, fake accounts, and other activities that undermine the positive and respectful interactions the LinkedIn community strives to uphold and keep them from impacting members, whether in the Feed or in private messages.

- **First Layer of Protection - Automatic prevention:** When a member attempts to create a piece of content on LinkedIn, various calls (or signals) are sent to LinkedIn's machine learning services. These services aim to automatically filter out certain policy violating content within 300 milliseconds of creation, meaning the content is visible only to the author and is not shown to anyone else on the platform.
- **Second Layer of Protection - Combination of automatic and human-led detection:** Detect content that is likely to be violative, but the algorithm is not confident enough about it to warrant automatic removal. Such content is flagged

by AI systems for further human review. If the human review team determines that the content violates LinkedIn's policies, it is removed from the platform.

- **Third Layer of Protection - Human-led detection:** Member-led protection involves LinkedIn members reporting the content on the platform. The content is then sent to LinkedIn's team of reviewers for further evaluation and is removed if found to be in violation of LinkedIn's policies.

Artificial intelligence (AI) is not new to LinkedIn. LinkedIn has long used AI to enhance its members' professional experiences. In 2025, an area of priority for LinkedIn was to strengthen its AI and Generative AI (GenAI) capabilities. Areas of focus included both the responsible application of GenAI to directly enhance LinkedIn user experience as well as increasing investments in the use of AI algorithms to further improve content moderation processes. Further leveraging the power of AI and expanding the use of GenAI features in the platform, supports LinkedIn's objectives to help its members connect, increase productivity, and achieve success in their careers.

While AI has enormous potential to expand access to opportunity and ultimately transform the world of work in positive ways, LinkedIn acknowledges that the use of AI comes with unique risks. That's why, consistent with its commitment to build a trustworthy platform, LinkedIn is committed to using AI responsibly. Inspired by, and aligned with, Microsoft's leadership in Responsible AI, LinkedIn uses the following Responsible AI Principles to guide its work:

- **Advance Economic Opportunity:** People are at the center of what LinkedIn does. AI is a tool to further LinkedIn's vision, empowering LinkedIn members and augmenting their success and productivity.
- **Uphold Trust:** LinkedIn's commitments to privacy, security, and safety guide its use of AI. LinkedIn takes meaningful steps to reduce the potential risks of AI.
- **Promote Fairness and Inclusion:** LinkedIn works to ensure that its use of AI benefits all members fairly, without causing or amplifying unfair bias.
- **Provide Transparency:** Understanding of AI starts with transparency. LinkedIn seeks to explain in clear and simple ways how its use of AI impacts people.
- **Embrace Accountability:** LinkedIn deploys robust AI governance, including assessing and addressing potential harms and fitness for purpose, and ensuring human oversight and accountability. LinkedIn is committed to learning from, and helping, others as AI best practices, norms, and laws evolve.

LinkedIn continues to advance its Trust and Safety efforts aligned with these principles and its commitments to advancing responsible innovation.

4. Products at a Glance

LinkedIn has members in over 200 countries and territories worldwide and operates a diversified business with revenue coming from its Talent Solutions (LTS), Marketing Solutions (LMS), Sales Solutions, and Premium Subscriptions products. While the Systemic Risk Assessment evaluates risk stemming from use or misuse of the LinkedIn online platform holistically, key products and features are described here to facilitate understanding of the platform's functioning and use.

4.1 Overall Product Design

LinkedIn adopts a safety-by-design approach across all its products, embedding trust and safety principles throughout its product development lifecycle. A dedicated Digital Safety team conducts pre-launch product reviews and collaborates with engineering teams to mitigate risks before launch. This includes safety features like fake profile detection and robust content moderation systems.

As part of this risk assessment, LinkedIn evaluated various ways in which the design of an online platform could accelerate or exacerbate harm to users to determine whether they were applicable to LinkedIn. In identifying the risk manifestations (i.e., the ways that risks may present on each LinkedIn features) for each assessed risk area, LinkedIn incorporated numerous ways that harms could manifest on each type of LinkedIn feature absent mitigation measures within the assessed risk manifestations. LinkedIn also considered a range of design choices, including features and platform behaviors believed to be linked to addictive behaviors. In general, based on quantitative and qualitative inputs, the assessment found that the LinkedIn platform is not designed in such a way that fosters addiction or otherwise exacerbates harm to users. LinkedIn's design emphasizes purposeful engagement rather than compulsive use. LinkedIn's recommender systems and AI tools are evaluated to avoid manipulative or polarizing content, supporting professional utility over addictive behaviors.


More specifically, LinkedIn's recommender systems are designed to recommend content based on quality over virality; therefore, LinkedIn's recommender systems generally do not contribute to the accelerated distribution of harmful content. Additionally, the professional and work-focused nature of the content on the platform significantly reduces potential addictive usage of the platform, as indicated by internally tracked usage data. The professional nature of the platform and the content shared and recommended on LinkedIn also has implications for whether specific design features may present harm to users. For example, while LinkedIn does allow leaderboards to show company scores for LinkedIn games, this professional application of leaderboards is not prone to the same social status comparison concerns that a leaderboard related to popularity or beauty-focused posts on a typical social media platform might. Furthermore, LinkedIn empowers users through transparent enforcement mechanisms, allowing them to report abuse and appeal moderation decisions. In addition, LinkedIn

empowers users through transparent enforcement mechanisms, allowing them to report abuse and appeal moderation decisions.

4.2 Existing Products

Profile: Each member account has a LinkedIn “Profile” associated with it. Functioning as an online resumé or curriculum vitae, the LinkedIn Profile serves as the member’s professional landing page on which they can manage their own personal, professional brand. Members must use their real or preferred professional names and include accurate information within their Profiles. Additionally, members must not misrepresent their identity, create a member Profile for anyone other than themselves (a real person), or use or attempt to use another’s account.


Figure 1: Profile



The Profile includes space for a personal headshot photo, as well as a “background” photo. There is an “about” section, as well as space for education history, work history, location, industry, contact information, skills, endorsements, and licenses and certifications, among other things. There is a “recent activity” section that shows snippets and thumbnails of recent posts or other actions by the member.

Using their Profile settings, members can control the extent to which certain aspects of their Profile information are accessible via LinkedIn’s in-product search feature (e.g., searching for a person by name or category), as well as by general browsing of the platform. For example, settings control who can see the member’s Profile photo and last name. Additionally, although portions of Profiles may be accessible by default via search engines to non-members or members in a logged-out state, members generally can control this via settings like those displayed in Figure 2.

Figure 2: Profile Public Visibility Settings




Feed: LinkedIn focuses on providing its members professionally relevant content through their Feed, which enables them to engage with people and topics they value. The LinkedIn Feed is the place on the platform where members go to view content, most of which is created by other members. Content can appear in a member's Feed because (1) the member follows or is otherwise connected through LinkedIn with the person or Page that posted it; (2) one of the member's LinkedIn connections reacted to, commented, or shared the content; or (3) it bears relevance to that member's job or industry or interests and/or is high quality professionally relevant content. Members may also see content from Groups they've joined, newsletters they follow, or Events they're attending. Members can customize and manage the content that appears in their Feed by following, unfollowing, blocking, or muting people, companies, and hashtags; hiding content in their Feed that they don't want to see; and reporting content that may violate LinkedIn's Professional Community Policies.

The following types of content can appear in the Feed:

- **Text:** This can be either short snippets / messages or long form posts that read more like an article or essay.
- **Images:** These can accompany text posts or be standalone images.
- **Video:** The Feed supports recorded videos, as well as links / embeddings of videos posted elsewhere. It is also possible to conduct a "livestream" event that appears in members' Feeds.
- **Ads:** Advertisements can appear in, or adjacent to, the Feed.
- **Notifications and Suggestions:** Suggestions for other people the member may want to connect with or follow, Pages and Groups the member may want to follow, LinkedIn Learning courses the member may want to take, and jobs the member may want to apply to may all appear in the Feed, as well as notices about connection updates or work anniversaries (subject to user settings).


A member's Feed experience is dynamic and specific to that logged-in member, based on, among other things, the member's LinkedIn Profile and network. Members can choose to view their Feed in a personalized or chronological manner. Logged-out users or guests do not have access to LinkedIn Feed.

Figure 3: Feed




Pages: Individual companies, organizations, and institutions can create their entity Profiles on LinkedIn (called LinkedIn Pages) to allow LinkedIn members to discover, follow, and engage with those organizations and learn more about them. To create LinkedIn Pages, LinkedIn members must verify that they are authorized representatives of their organization and agree to the LinkedIn Pages Terms. Organizations with LinkedIn Pages can have one or more LinkedIn member Page administrators who manage the Page and post content (such as industry updates, information about the organization, and job posts) to build and promote the organization’s brand to the professional audience on LinkedIn.

Figure 4: Page




Groups: LinkedIn Groups are forums members join to discuss shared interests or experiences. Groups function similarly to Pages, but unlike Pages, Groups are created and maintained by individual members rather than organizations or institutions. A Group can be public or private. Group members can send message requests to other Group members and start discussions or conversations in the Group by posting in the Group's Feed text, images, a video, or a link. Group members may see each other's Group posts, likes, and comments within the Group, or in their Feed.

Figure 5: Group




Newsletters: LinkedIn's Newsletter feature allows logged-in members to author a set of regularly published articles on LinkedIn, typically about a specific professional topic that interests the author. Each LinkedIn Newsletter has its own feed where members can learn about the newsletter and its author, see past editions, and subscribe to or share the Newsletter. This feed includes an overview of the Newsletter, the author's Profile information, and a list of articles in the newsletter.

Figure 6: Newsletter




Messaging: LinkedIn’s messaging feature allows members to contact and communicate with other LinkedIn members through direct messages. LinkedIn members can send (1) a message to a connection or multiple connections in a Group chat; (2) a message request to certain non-connections based on a shared context (i.e., a message request to a fellow Group member or coworker at your current company); or (3) a connection request with a personal message to other members. In addition, “InMail” messages are a premium (for a fee) feature that requires the use of credits to message non-connections.

Figure 7: Messaging



Jobs: The LinkedIn Jobs page contains job listings posted by employers or compiled from related sites and services. There are various options to apply for Jobs via LinkedIn, though some listings only allow applicants to apply for a Job directly from the employer’s own application website. Recommended Jobs are curated for the individual member based on their skills, job titles, education, interests, and various other attributes, including Jobs that the member has specifically indicated to LinkedIn they are interested in. Members are also able to search for Jobs by title, company, and other attributes.

Figure 8: Job




Search: Using LinkedIn’s in-product search feature located at the top of most LinkedIn pages, LinkedIn members can conduct a search across most of the platform, including posts and Profiles. A searching member has the option to “See all results” for a keyword search or to narrow a search by categories such as Jobs or people.

Figure 9: Search



Events: LinkedIn Events are an easy way to create and join professional events, such as meetups, online workshops, seminars, and more. Members can find and join communities, grow their business, network with others, and learn new skills. There are two types of events: (1) In-person events taking place offline in a physical venue and promoted on LinkedIn; and (2) Online events that are either (a) streamed on LinkedIn, with real-time conversations in the comments and a recording on a Page or Profile; or (b) hosted off LinkedIn and promoted on LinkedIn.


Figure 10: Events



Games: LinkedIn offers daily, thinking-oriented games designed to allow users additional opportunities to engage within their networks and spark conversations. The following games are available to EU users: (1) Pinpoint — a word association game; (2) Crossclimb — a timed trivia game; (3) Queens—a logic game¹. Additionally, LinkedIn launched two new games since August 2024 which are discussed in detail in the [New Since Last Year section](#). Members who play one of LinkedIn’s games can share scores with their network, see their streaks, and find leaderboards. Games are available to all logged-in members. There are no prizes for winning games.

- **Pinpoint**—a word association game:

Figure 11: Pinpoint



¹ All LinkedIn games do not allow users to interact within partially or fully simulated virtual environments


- **Crossclimb**—a timed trivia game:

Figure 12: Crossclimb



- **Queens**—a logic game:

Figure 13: Queens




Video Tab: LinkedIn’s Video Tab is a feature shown at the bottom of certain LinkedIn mobile experiences. When a member clicks on the Video Tab icon, it leads them to an immersive video player that allows members to swipe between editorially curated short form videos. The Video Tab only features video content stemming from the Feed, and each video in the Video Tab features the same underlying post, comments, and reactions as the original video in the Feed.

Although the Video Tab was available to members in the European Union/European Economic Area during a portion of the risk assessment period, this feature currently is not available to EU-based members.

LinkedIn Marketing Solutions (LMS): LinkedIn’s advertising platform primarily consists of fully integrated first-party advertising on LinkedIn with a focus on “business-to-business” marketing (i.e., one business advertising to another business, as opposed to a consumer). For members in Europe, there are four types of ads on LinkedIn: Sponsored Content, Text Ads, Sponsored Messaging, and Dynamic Ads:


- **Sponsored Content:** These are native posts in the Feed from the advertising customer. They can feature text, images, a video, or a job posting and links to the advertiser’s website of choice. These posts are identified as advertisements by terms such as “Promoted” that appear in the upper-left corner of the post.

Figure 14: Sponsored Content




- **Text Ads:** These include a headline, brief text, and an image. They are shown either above or to the right of the Feed. These are also identified as advertisements by a “Promoted” label.

Figure 15: Text Ads




- **Sponsored Messaging:** These ads are served in the context of the LinkedIn Messages interface and include a text ad and link or call-to-action button. They are identified as advertisements with a “Sponsored” label in the message preview and the message itself. They appear as native ads in a member’s LinkedIn messaging inbox. EU-based members must provide affirmative consent to receive Sponsored Messaging ads. Otherwise, they may receive ads displayed in a rotating banner above their inbox.

Figure 16: Sponsored Messaging



- **Dynamic Ads:** These ads use the member's Profile image, name, or job function to create customized ads directed toward that member. Dynamic ads appear to the right of the Feed and are marked with an "Ad" label in their top-right corner.

Figure 17: Dynamic Ads



LinkedIn members may receive targeted ads based on information provided and inferred from their LinkedIn Profile and usage information. For example, years of experience can be inferred from the sum of the years of work history indicated on a member's Profile and professional interests can be inferred from LinkedIn Feed posts. Members can view details about why specific ads are displayed to them by clicking the three-dot icon in the upper-right corner of the ad and selecting "Why Am I Seeing This Ad?"

Using their [ad settings](#), a member can manage general advertising preferences, how data collected on LinkedIn is used to show the member ads, and how third-party data is used to show the member ads. If a member limits the information used to improve ad relevance, that member will still see ads on LinkedIn, but they may not be related to the member's interests and/or experience.

Verification: LinkedIn's voluntary verification features allow members to verify certain information about the Pages they administer, the Jobs they post, or their Profile (like the member's association with a particular company or educational institution, or their identity, using a government-issued ID through one of LinkedIn's identity verification partners). Verified information, which is marked by a badge symbol in various locations on the platform, provides members with authenticity signals about the people, Jobs and Pages they encounter on LinkedIn. These signals are designed to help members make more informed decisions about connecting or engaging with other professionals, organizations, or Jobs.


Currently, LinkedIn supports several forms of verification. Each form of verification has its own eligibility and availability criteria. Verifications are free, aimed at increasing trust and authenticity on the platform, and can be removed at any time by the member or entity whose information is verified.

Figure 18: Verification badge




- **Verifications on member Profiles:** As described in greater detail below, members can add an identity, workplace, or educational institution verification to their Profile. Once the member has successfully done so, a verification badge is displayed next to the member's name on their Profile and information about the verification is included in their "About this Profile" section. The badge may also appear next to the member's name elsewhere within LinkedIn products (e.g., next to Jobs the member posts). This information will be visible to logged-in members.

Figure 19: Profile Verification



- **Identity Verification:** EU Members may verify their identity through LinkedIn’s verification partner, Persona. LinkedIn uses different identity verification partners for other countries and regions, including CLEAR in the United States, Canada and Mexico and DigiLocker in India. As shown below, LinkedIn provides basic information about how the verification process works for members with Near-Field Communication or NFC-enabled government IDs (e.g., E-passports) issued in the EU. If a member opts to “Verify with Persona,” they are redirected to Persona, which will collect an image of the member’s government issued ID and a live selfie.

Figure 20: Identity Verification




If the member successfully completes Persona’s verification process, LinkedIn does not receive a copy of any government IDs from Persona, but -with the member’s consent- Persona shares certain limited data about the member with LinkedIn (such as the member’s name as it appears on the government ID). LinkedIn uses this data to add a verification to the

member's Profile. LinkedIn may also use this data for security purposes (e.g., to prevent identity theft and fraudulent use of LinkedIn's verification processes), in accordance with LinkedIn's Privacy Policy.


- **Workplace Verification:** Members may confirm their association with a specific company using their work email, Microsoft Entra Verified ID, or a LinkedIn Learning license or active LinkedIn Recruiter license provided by the company.
- **Verified Recruiter Label:** When a member (1) verifies their workplace using a LinkedIn Recruiter license and (2) currently lists the company that issued that license as their current employer in the "Work experience" section of their Profile, a "Verified Recruiter" label will appear near the member's name in InMails sent via LinkedIn's Recruiter product.

Figure 21: Recruiter Verification



- **Educational Institution Verification:** An educational institution verification indicates that the member has confirmed their association with an educational institution by using an email address containing a domain allowed by the institution, as indicated by the super admins of a verified institution's LinkedIn Page. Currently, the feature is only available to select educational institutions.
- **Page Verification:** LinkedIn will display a verification badge next to the organization name on LinkedIn Pages that meet certain criteria (e.g., active Page admin presence, compliance with LinkedIn's policies, and accuracy of certain data points on the Page, like company location and URL). This verification badge helps signal that the Page is an official Page of the organization, which helps organizations build credibility with potential customers and helps members find the actual organizations they intend to engage with.

Figure 22: Page Verification



Once a Page is verified: (1) LinkedIn members can verify their employment with the Page's organization using a Microsoft Entra Verified ID (if any), or a LinkedIn Learning or LinkedIn Recruiter license provided by the organization; and (2) the Page's super admins can enable members to verify their association with the organization using an email address containing domains controlled by the super admins.

- **Verifications on Jobs:** To provide job seekers additional trust and authenticity signals as they apply for Jobs posted on LinkedIn, LinkedIn displays a verification badge on Job posts. When members see a verification badge on a Job post, it means that certain information about the company has been confirmed by LinkedIn, or information about the job poster has been verified by the job poster through LinkedIn or through a trusted third-party partner.

GenAI Features: LinkedIn's GenAI features are integrated into existing products to enhance LinkedIn member experiences, providing members with options to more easily consume and share information. LinkedIn's GenAI features fall into two categories (1) the generation of new content and (2) the generation of insights based on existing content. These features include:

Generation of new content:

- **Feed:** LinkedIn's [Writing Assistant](#) can create draft LinkedIn posts (i.e., draft textual content which a member may choose to post on the platform) based on Premium subscriber prompts.
- **Messaging:** LinkedIn's Writing Assistant can help members create a first draft message (text only). Premium members can create a draft message using pre-set prompt suggestions or by typing in their own prompt. Recruiters can also create draft messages (for potential job candidates).
- **LinkedIn Talent Solutions (LTS):** For job posters, LinkedIn's AI-Powered job descriptions help recruiters create job postings with initial information provided by the job poster, such as job title, company name, workplace type, job type, and location, or from the job poster's profile, such as company and location. Recruiters can review and edit all elements of the draft job posting (including the AI-suggested text) prior to publication.
- **LinkedIn Marketing Solutions (LMS):** LMS GenAI features are integrated into LinkedIn's existing Campaign Manager advertising platform and designed to simplify advertisement development and campaign initiation. Campaign Manager incorporates Microsoft Designer, a third-party design tool that leverages GenAI to empower marketers to create unique and high-quality graphics via prompts, in less time, and without any design background. Advertisers can also leverage AI to create a draft subject line for their conversational advertisement, which is an advertisement within the LinkedIn messaging UI.

- **Collaborative Articles:** Collaborative articles are textual articles focused on a specific professional topic that are published by LinkedIn with insights and perspectives contributed by LinkedIn members. Articles begin as AI-powered conversation starters developed with oversight from a LinkedIn team.
- **Ancillary Features:** Additional ancillary GenAI features are incorporated throughout LinkedIn to enhance the value of the platform to its members. For example, Virtual Chat Assistant (VCA) is a text-based interface that offers AI-generated suggested responses to customer questions, drawing primarily on the help center content already available to those customers. For unresolved queries, the customer is connected with a live agent.

Generation of insights based on existing content:

- **Profile:** LinkedIn's AI-powered Writing Assistant shares personalized suggestions upon request to enhance the Headline and About section of a Premium subscriber's Profile. The Resume to Profile Suggestions feature also allows members to easily transfer information from an uploaded resume to their LinkedIn Profile. Members can also use Post Ideas to suggest questions or topics a member could write about in their next post based on the experience included in their Profile.
- **Feed:** AI-powered Insights on LinkedIn use GenAI to provide members with personalized insights and key takeaways on posted content.
- **LinkedIn Talent Solutions (LTS):** For job seekers, AI-powered job insights will help job seekers learn more about a job or company, understand their fit for a role, and how to best position themselves if they apply. This includes recommending relevant jobs and providing interview tips. Premium subscribers can also use GenAI to tailor their resumes and cover letters for a particular role. For enterprise seat holders of LinkedIn Recruiter, AI-Assisted Search and Projects enable recruiters to quickly search for candidates, create projects, and get recommendations by typing out needs into prompt box and chat experience.
- **LinkedIn Marketing Solutions (LMS):** Accelerate is LinkedIn's AI-powered advertising campaign type, which can help advertisers to more easily create campaigns. Accelerate aims to improve return on ad spend by continuously finding the right combination of targeting, creative, bidding, and placement for ad campaigns. When an advertiser creates an Accelerate campaign, they will provide information about the product, service, or brand, which is then analyzed to automatically generate suggested budget, schedule, audience, creative, and ad placements, each of which is reviewed by the advertiser before launching.

4.3 New since Last Year


LinkedIn's approach to product updates: At LinkedIn, new products and initiatives undergo a pre-launch review [process](#) to identify and address potential associated risks and abuse vectors that are likely to impact systemic risk. In cases where that impact is potentially critical, LinkedIn may initiate an out-of-cycle assessment of systemic risk to ensure adequate mitigations are in place prior to the launch. While no product updates this year met the criteria for an out-of-cycle update to the Systemic Risk Assessment (no updates were determined to be likely to have a critical impact on the risks identified pursuant to DSA Article 34), highlighted below are several changes to the service that were reviewed and implemented since the last Systemic Risk Assessment report.

LinkedIn Hiring Assistant: In October 2024, LinkedIn introduced its first AI agent, Hiring Assistant within LinkedIn Recruiter. This tool is designed to handle a recruiter's most repetitive tasks, allowing them to focus on more impactful work like advising hiring managers, connecting with candidates, and creating positive candidate experiences.

Hiring Assistant, an add-on to the LinkedIn Recruiter platform, helps recruiters to more efficiently and effectively identify qualified candidates for a role. Recruiters can provide a job description, an ideal candidate's name, a LinkedIn job post URL, or upload a file, and Hiring Assistant will generate the project or job based on the provided information. It will recommend both required and nice-to-have qualifications for the project, source candidates, and streamline the applicant review process.


Additionally, Hiring Assistant helps recruiters in reviewing potential candidates' qualifications. When evaluating candidates with Hiring Assistant, each candidate profile includes a summary explaining why the Hiring Assistant believes they are a good fit and how many qualifications they meet. This assessment is based on information from sections of the candidate's LinkedIn profile, recruiting notes, and the candidate's resume and responses to screening questions, to the extent that these are complete or available.

Figure 23: Hiring Assistant Top Applicants



Hiring Assistant also enables recruiters to send AI-generated messages to up to 25 candidates at once. Once the Recruiter messages candidates, a pre-screening agent can help the Recruiter by collecting additional information from candidates.


Figure 24: Assisted Pre-screening



Hiring Assistant can be enabled or disabled on any project that recruiters own. LinkedIn also introduced Hiring Assistant features into Online Jobs to help job posters find, source and review candidates using the same underlying AI models as Hiring Assistant for Recruiter.

LinkedIn Sales Assistant: During the assessment period, LinkedIn introduced an AI agent called Sales Assistant within its Sales Navigator product. This Sales Assistant empowers Sales Navigator customers to perform AI-enhanced search queries, utilize AI-driven company insight summaries, generate talking points for engaging potential sales leads on LinkedIn, draft effective messages to leads and customers more efficiently, and prepare for meetings by providing AI-generated insights from Sales Navigator in meeting invitations.

Figure 25: LinkedIn Sales Assistant




Enhanced member education and use of feature restrictions in content moderation: LinkedIn operates under standards of professionalism, which are reflected in its content policies. These policies bolster a safe, trusted, and professional platform, and LinkedIn strictly enforces them. Where content violates LinkedIn's [Professional Community Policies](#), LinkedIn may limit its visibility, label it, or remove it entirely, depending on the severity of the violation. Historically, repeated violations of the Professional Community Policies also resulted in the restriction of the member's LinkedIn account. Account restriction has been a key measure to prevent the ongoing posting of harmful content on its platform.

LinkedIn's services are important to members' economic opportunity. Therefore, account restriction — terminating a member's access to LinkedIn's services entirely — can be a severe penalty. Restricted accounts are not viewable on the platform (including by recruiters) and cannot view or engage with content on LinkedIn or search for and apply for jobs. To better balance the prevention of harmful content with professional expression, during the assessment period LinkedIn began implementing feature restrictions instead of account restrictions under many of its policies. To increase member reform opportunities, LinkedIn also enhanced its in-product member

education about its policies. These changes, as described below, aim to take a more tailored and proportional approach to mitigating harms and other policy violating activity as they occur on LinkedIn, balancing the harm of violative content with the harm of limiting professional expression.

- **Member Education:** Where LinkedIn actions policy-violating content, in addition to sending its existing content moderation notices and redress information (i.e., LinkedIn's Article 17 Statements of Reasons), LinkedIn also now surfaces in-product policy reminders to the member whose content was actioned. These reminders appear in-product when the member attempts to post, comment, or otherwise share additional user-generated content through LinkedIn. They provide a link to the policy that the member violated and a reminder to keep additional conversations on LinkedIn professional and compliant with the Professional Community Policies. These reminders are designed to better educate members about what is not permitted on LinkedIn and to encourage professional engagement at the time members are deciding what to share on platform.


Figure 26: In-product Policy Reminders



- **Feature Restriction(s):** If, after viewing the in-product notices described above and receiving separate Statements of Reasons, a member repeatedly violates LinkedIn's Professional Community Policies, LinkedIn will now restrict the member's access to certain features, like creating content (including posting, reposting, or editing past posts), commenting or editing past comments, or starting a new message conversation with connections of any degree. The member will remain able to view content, apply for jobs, reply to recruiters or anyone who directly messages the member, continue existing private conversations, delete the member's past posts or comments, block other members, and control who can see or comment on the member's existing posts.


LinkedIn's Statements of Reasons to members make clear that repeated policy violations may result in permanent feature restrictions.

Figure 27: Statement of Reasons



LinkedIn notifies the member when it applies a feature restriction to their account and provides the member information about how to challenge the feature restriction (namely, by appealing the underlying moderation decisions that triggered it). For safety reasons, LinkedIn will continue to restrict member accounts for certain types of violations, like those involving child sexual abuse material and terrorism.

Figure 28: Feature Restriction Notices




Brandlink: During the assessment period, LinkedIn launched a beta advertising program, Brandlink, which allows certain brands to connect with trusted publishers and creators to promote in-stream, pre-roll video ads alongside contextually-aligned, high-quality videos directly in targeted members' LinkedIn newsfeeds.

Games: Due to member interest in LinkedIn games, LinkedIn launched two new daily thinking-oriented games, Zip and Tango, to provide members additional opportunities to engage within their networks and spark conversations. Zip is a logic game where the goal is to draw a path through the grid by connecting the numbers in order while making sure every cell is used and. Tango is once-a-day logic puzzle for which the goal is to fill the grid so that each cell contains either a “sun” or a “moon” without placing more than two of the same next to each other.


Zip — a logic game

Figure 29: Zip



Tango — logic puzzle:

Figure 30: Tango



5. Methodology

LinkedIn aligned its risk assessment process and methodology with the DTSP's Safe Framework, the industry standard for the evaluation of online safety best practices. LinkedIn also engaged a global consulting firm with experience and expertise in the conduct of risk assessments across industries, including technology, energy, climate, finance, and human rights, to support the development of the process and methodology. During the assessment, probability of risk manifestation and severity were considered as inputs to assess the Inherent Risk (risk absent mitigations). Reasonableness, proportionality, and effectiveness of mitigations (informed by the best available information and scientific insights) were also considered as inputs for scoring and factored into calculating the Residual Risk (risk after mitigation).

LinkedIn followed the Safe Framework process to conduct a full assessment including **Discovery, Identification, Assessment, Testing, and Reporting**, which it believes to be proportionate not only to the platform but to the risks identified as well.

Throughout the process, LinkedIn examined risk across the twelve risk areas identified in Article 34. The methodology is summarized here and further expanded in [Appendix A: Methodology Expanded](#).

5.1 Year 3 Design Enhancements


This year, LinkedIn followed the same methodology used in the Year 2 Systemic Risk Assessment, with some enhancements to incorporate lessons learned and enable a more tailored consideration of specific systemic risks to relevant mitigations. These enhancements include the following:

- Separating certain Risk Scenarios within Risk Areas into multiple discrete Risk Scenarios to more explicitly identify which risks are considered under each risk area, incorporating both risks specific to DSA Article 34 as well as LinkedIn's more holistic understanding of potential risks to users.
- Shifting toward a more data-driven maturity of mitigations calculation, by identifying gaps in coverage of risks, calculating the depth of coverage of each risk, and calculating the breadth of mitigation coverage across the DTSP Best Practice framework.
- Additions to this report to incorporate certain external general feedback civil society has provided to industry, including for example to explain more fully how LinkedIn considers the ways in which the design of its platform, services, and features could contribute to or mitigate certain harms.

5.2 Risk Assessment Process

LinkedIn adopted and tailored the Safe Assessment Framework to perform its Systemic Risk Assessment.

Figure 31: LinkedIn's Risk Assessment Process



Discover

As part of discovery, LinkedIn collected responses to the previously modified DTSP Safe Assessment questionnaire (which gathers information on best practices implemented to mitigate risks); examined external authoritative sources and guidance (such as guidance and feedback from the European Commission and civil society organizations on systemic risk assessments); and collected and reviewed relevant public policies and publications to identify any adjustments needed to the Systemic Risk Assessment process this year.

Inputs to Risk Assessment:

- DTSP questionnaire responses
- Authoritative sources
- Policies and publications

Identify

LinkedIn built upon last year's risk identification process and focused this year on identifying updates and changes in risk trends, the product itself, and implemented mitigations. In support of this process, LinkedIn collected written updates from stakeholders on the manifestation and mitigation of risks across products and Risk Areas; reviewed relevant internal controls; conducted stakeholder engagements to dive deeper into products with relevant feature updates this year; conducted a tailored engagement with a civil society organization; and analyzed key internal and external metrics related to Trust and Safety policy enforcement as well as trends in potential systemic risk related to LinkedIn.

Inputs to Risk Assessment:

- Product and feature changes
- Internal consultations
- External consultations
- Internal content moderation metrics
- Transparency reporting metrics
- Open-source data on public discourse related to LinkedIn and systemic risk areas
- DSA Risk and Control Matrix

Assess

LinkedIn assessed the probability and severity of each of the twelve systemic risks stemming from use or misuse of its platform and services, as well as the maturity of risk mitigation measures implemented to arrive at a prioritization of systemic risk areas with potentially higher levels of Residual Risk. The assessment incorporated the referenced ten inputs from the Discovery and Identification phases to inform scores and implemented a traditional risk assessment equation to aid in prioritization of risk for action.

LinkedIn employed a **data-driven probability assessment**, combining insights from internal, externalized, and open-source data to assess the likelihood of specific risks stemming from use or misuse of LinkedIn considering the vulnerability of the platform *absent* mitigation measures and user demand or likely frequency of attempts to perpetrate each harm on the platform. This methodology follows models used in climate and energy impact assessments.

LinkedIn leveraged an **objective, systems-based assessment of severity**, examining the systems impacted² by the systemic risk area (environmental, geographic, political, security, societal, and wellbeing), including scale and gravity of impact (for example, fraud poses a high economic risk at the individual level, but a lower environmental risk at the global level). This methodology follows models used in environmental impact assessments.

² **Environmental System** is the interplay of natural processes and ecosystems that sustain life on Earth, including the atmosphere, land, water, and biodiversity; **Geographic System** is the physical and spatial environments that shape natural landscapes and human societies including the distribution of climates and resources across the Earth; **Political System** is the structures, institutions, and processes through which decisions are made and power is allocated within a society, shaping governance and public policy; **Security System** is the measures and institutions dedicated to protecting a society from internal and external threats, ensuring the safety and stability of the state and its citizens; **Societal System** is the social relationships, cultural norms, and institutions that bind individuals together, influencing values, behaviors, and community life; **Wellbeing System** is the overall health, happiness, and prosperity of individuals, encompassing physical and mental health, economic security, and quality of life.

LinkedIn aligned existing mitigations and controls to the 35 DTSP Best Practices and evaluated their implementation according to the **DTSP Maturity Rating**, while also considering their proportionality, reasonableness, and effectiveness relative to the potential manifestation of risks across the LinkedIn product. This methodology follows the industry standard for evaluation of Trust and Safety best practices.

LinkedIn also considered for each applicable Risk Area and mitigation whether and how the following factors influence any of the systemic risks: the design of recommender systems and other relevant algorithmic systems; content moderation systems; applicable terms and conditions and their enforcement; systems for selecting and presenting advertisements; data-related practices; intentional manipulation including inauthentic use or automated exploitation; amplification and potentially rapid and wide dissemination of illegal or violative content; and linguistic and regional considerations.

LinkedIn considered the Inherent Risk for each risk area resulting from combined ratings of probability and severity and reduced the Inherent Risk rating by a percentage proportional to the assessed strength of relevant mitigations. This methodology produced a view of Residual Risk across the risk areas to enable prioritization of specific risk areas for further mitigation, in alignment with Article 35 of the DSA.

Test

Quality Assurance is the process of taking systematic efforts to ensure that the product delivered to LinkedIn's members meets with their performance, design, reliability, and maintainability expectations. LinkedIn routinely conducts quality assurance activities and tests the effectiveness of its algorithmic systems and content moderation processes. LinkedIn has also implemented a more quantitative assessment of the maturity of mitigation measures in addressing systemic risks as part of the annual assessment process. The DSA Compliance Program monitors the effectiveness of systemic risk mitigations through formal controls, such as the monitoring of prevalence metrics and uses the results of the third-party audit of DSA compliance as an additional validation of these controls.

Report

This Report fulfills step 5 of the DTSP Safe Assessment Framework process in addition to Article 42(4) of the DSA. LinkedIn reserves the right, as indicated in DSA Article 42(5) to remove such information from the publicly available version of this Report that may result in the disclosure of confidential information, cause significant vulnerabilities for the security of the service, or undermine public security or harm recipients.

5.3 Risk Assessment Tools

LinkedIn employed a Systemic Risk Assessment workbook to facilitate the standardized completion of Systemic Risk Assessment and to store relevant ratings, scores, and rationale. The Risk Assessment team used this workbook to collect and assess the relevant risk factors and risk scenarios, to identify potential manifestation of risks across LinkedIn products and services absent safeguards, to capture and align implemented mitigations, to map mitigations to identified risks, and to calculate and display all risk assessment scores.

6. Summary of Results

LinkedIn assessed systemic risk stemming from the functioning, use, or misuse of the platform and services within twelve key risk areas. To identify the residual systemic risk rating across each Risk Area, LinkedIn identified respective mitigations that actively work to address the Risk Scenarios in one or more Risk Area. Where a mitigation applies to more than one Risk Area, it may be discussed in more than one location throughout this Report. The full Risk Areas, definitions, and Risk Scenarios are included in [Appendix A: Methodology Expanded](#).

6.1 Key Findings

Risk Scores

Using the above methodology, LinkedIn calculated the risk scores listed below for each of the twelve Risk Areas covered by this Report. Definitions for each potential Inherent Probability, Inherent Severity, Inherent Risk, Mitigation Maturity, and Residual Risk rating can be found in [Appendix A: Methodology Expanded](#).

Figure 32: 2025 LinkedIn Systemic Risk Assessment Scores

Risk Area	Inherent Probability	Inherent Severity	Inherent Risk Rating	Mitigation Maturity	Residual Risk Rating
Consumer Protection and Fraud	Expected	High	Critical	Managed	Moderate
Human Dignity	Likely	Critical	High	Defined	Moderate
Mental and Physical Wellbeing	Likely	Critical	High	Managed	Low
Discrimination and Hate	Likely	High	High	Managed	Low
Freedom of Expression and Information	Likely	High	High	Managed	Low
Protection of Personal Data	Likely	High	High	Managed	Low
Civic Discourse and Electoral Processes	Not Likely	Critical	Moderate	Managed	Low
Public Health	Not Likely	Critical	Moderate	Defined	Low
Illegal Content and Activities	Not Likely	Critical	Moderate	Managed	Low
Private and Family Life	Not Likely	High	Moderate	Managed	Low
Public Security	Remote	Critical	Low	Managed	Minimal
Rights and Protection of Minors	Remote	Critical	Low	Managed	Minimal

Risk Areas of Moderate Residual Risk

Considering the probability, severity, and maturity of mitigations, LinkedIn assessed that two of the assessed risk areas fell into a Residual Risk category of Moderate:

Consumer Protection and Fraud and Human Dignity.

Consumer Protection and Fraud

Risks related to Consumer Protection and Fraud include scamming, defrauding, spamming, or falsifying information, including the creation of fake accounts and use of bots or other deceptive commercial practices on LinkedIn. Absent sufficient mitigations, risks related to Consumer Protection and Fraud may manifest on the platform in ways such as users creating fake profiles and defrauding members through private messages or posting scams on Feed, advertisers posting links to fraudulent content or malware, or users posting fraudulent job postings.

The **Inherent Probability** of systemic risks concerning Consumer Protection and Fraud stemming from the function, use, or misuse of LinkedIn's services is assessed as "Expected." The LinkedIn platform favors knowledge sharing rather than virality, and therefore, tends to not facilitate the rapid and wide dissemination of content that does not meet a minimum quality bar. In turn, this limits the systemic risk posed by the spread of harmful content. However, while there is a lower likelihood of members sharing fraudulent content on LinkedIn due to its professional nature, there is a higher risk of bad actors creating fraudulent accounts and using them to defraud members due to the financial incentives of fraudulent activities. Out of the twelve risk areas, Consumer Protection and Fraud received the highest probability score due to the relative frequency of content or activities related to Consumer Protection and Fraud appearing in LinkedIn's content moderation systems as well as public discourse related to LinkedIn. Therefore, the probability is assessed as "Expected" that users could share or be exposed to fraudulent content or activities negatively impacting Consumer Protection on LinkedIn's platform absent sufficient mitigations.

The **Inherent Severity** of impact for fraudulent content or activities on the LinkedIn platform is rated as "High" due to the gravity of risks within this category and the potential for significant economic impact at the individual level. While economic, and even security, societal, and wellbeing, impact at the individual level can be severe, the risk is generally limited to the individual level and does not necessarily scale to the country, regional, or global impact scale. Not all risks within this category would be considered "High," as the impact of spam and deceptive advertisements are less likely to cause significant harm, compared to malware links or deceptive Job postings.

With the rating for Inherent Probability as "Expected" and the Inherent Severity as "High," the **Inherent Risk** associated with negative impacts to Consumer Protection and Fraud across LinkedIn products and features is therefore assessed as "Critical."

In addition to the DTSP Best Practices LinkedIn has effectively implemented across all Risk Areas, which are further described under “[Mitigations](#),” LinkedIn has implemented robust mitigations specific to Consumer Protection and Fraud following industry best practices with defined, documented, and managed processes. As such, the **maturity of mitigation** efforts LinkedIn has applied to the Consumer Protection and Fraud risk area is assessed as “Managed,” which brings the **Residual Risk** rating down to “Moderate.”

Considering the rapidly evolving nature of threats related to Consumer Protection and Fraud and the need to continually assess and enhance detection techniques, LinkedIn continues to prioritize the development and enhancement of strategies to further manage and mitigate risks related to Consumer Protection and Fraud, aiming to maintain a high level of consumer safety on the platform. The key mitigations currently implemented specific to Consumer Protection and Fraud are described in the subsequent paragraphs of this section.


LinkedIn prohibits Profiles with false information including fake identities, misleading or deceptive information regarding business, qualifications, work experience, affiliations, or achievements. Members must use a Profile photo that accurately reflects their likeness and not associate with businesses that are not relevant to their professional associations. Creating or using another member’s account or Profile is not allowed, and sharing LinkedIn accounts is prohibited. Scamming, defrauding, or deceiving others on the platform are strictly prohibited, as is promoting pyramid schemes or romance scams, sharing malicious software, or engaging in phishing attempts. LinkedIn has a dedicated abuse prevention team that prioritizes, triages, and remediates numerous abuse vectors (both acute and systemic) in order to keep LinkedIn members safe. LinkedIn’s Trust Product team orchestrates regular prevalence exercises where it manually reviews samples of accounts and content on platform to identify new patterns of abuse on the site, which inform both design of future defenses and priorities in its Trust and Safety roadmap. The Trust Product team has also implemented several proactive defenses to detect fraudulent behavior and policy-violating content. LinkedIn monitors the performance of these defenses regularly to ensure they are working optimally and address any gaps as they are identified.

LinkedIn also has a dedicated Account Abuse team that works to action LinkedIn’s impersonation defense strategy, which encompasses verifying Profiles, performing a name to title to company match, tracking of member IP addresses, and leveraging member feedback to identify fake accounts at scale. LinkedIn applies a multi-layer approach to detecting fake accounts on the platform.

- First, to proactively stop fake accounts at scale during registration, LinkedIn uses machine learning models to detect groups of accounts that look or act similarly. Every new user registration attempt is evaluated by a machine-learned model that gives an abuse risk score. Signup attempts with a low abuse risk score are allowed to register right away, while attempts with a high abuse risk score are

- prevented from creating an account. Attempts with medium risk scores are challenged by LinkedIn's security measures to verify that they are real people.
- Second, LinkedIn creates clusters of accounts by grouping them based on shared attributes. LinkedIn then finds account clusters that show a statistically abnormal distribution, scores the clusters with supervised machine learning models, and then actions accounts in those clusters.
- Third, LinkedIn looks for activity-based models that either look for specific types of bad behavior typical of abusive accounts (e.g., excessive connection or invitations sent) or behavior that is anomalous.
- Fourth, LinkedIn gets signals from its members. Members give valuable information by reporting previously undetected accounts that can then go through additional model scoring and review.
- Finally, LinkedIn manually reviews accounts based on internal findings, peer leads, and member flags.

Figure 33: Fake Account Defense Funnel



LinkedIn has also implemented a variety of verification features. To help these in-product authenticity signals serve their purpose (i.e., as user empowerment tools to inform user engagement on the platform), LinkedIn clearly marks verifications associated with Profiles, Pages, or Jobs with a verification badge icon and corresponding transparency about what such verifications means. This design and transparency aims to reduce the risk that users overly rely on or otherwise misinterpret verifications. LinkedIn also takes various steps to help prevent abuse of LinkedIn's verification features. For example, in the EU, members who verify their identity through LinkedIn's partner Persona must use certain hard-to-falsify documents like E-passports. In jurisdictions where it is legally allowed, LinkedIn's identity verification partners independently use liveness checks to confirm that the picture in the ID document submitted to the partner matches the member's face.

LinkedIn also offers several other features to protect its members from inauthentic accounts. Members can choose to only receive communications from people they know, thus allowing them to control who can contact them. The “About this Profile” section of a LinkedIn Profile shows members when a Profile was created (i.e., when the member joined LinkedIn) and if the Profile has any verifications associated with it (like a workplace or identity verification). For messaging, members can further protect themselves by opting into settings that detect harmful content in messaging (e.g., crypto scams), as well as limit the people who can attempt to connect with or message them. Once enabled through enhanced safety features in member account settings, LinkedIn also displays safety reminders to members if attempts to move a private conversation away from the platform are detected in their incoming messages. The member can click through to view or report the message. This flow is designed to empower and educate members to make informed safety choices before they move certain conversations off platform. For Jobs, LinkedIn has a range of mechanisms in place to help protect job seekers from fraud and scams.

- LinkedIn has implemented anti-abuse models and processes for machines and humans to work to identify, review, and restrict fraudulent job posts.
- LinkedIn has implemented systems to provide job seeker safety tips to job seekers directly in-product on LinkedIn’s platform, before they apply to a job, when they are searching for a job, and after they apply to a job, as appropriate.
- LinkedIn also has implemented in-product experiences and provides resources to remind members that while they are in their job search, they should keep an eye out for signs of potential fraud including jobs that ask for payment, account credentials, or other highly sensitive personal data; that claim to offer high pay for little work; that give job offers after one remote interview; that include bad grammar, multiple misspellings, and grammatical errors; that ask you to purchase your own equipment (computers, iPads, mobile phones) and request you to mail it to their IT department; and so on. LinkedIn also reminds members to research a company, its official website, and its social media accounts before applying, and encourages members to report suspicious jobs to LinkedIn.
- LinkedIn has implemented systems to allow users to report suspicious, fraudulent, and policy-violating job posts to LinkedIn. From each job post, members can report the job post as a “scam, phishing, or malware.”
- LinkedIn has implemented systems to help companies protect their LinkedIn Pages from fraudulent or unauthorized jobs posted by bad actors. Job posts on LinkedIn must be associated with a company’s LinkedIn page. LinkedIn provides companies with the ability to protect and limit who can post jobs to their LinkedIn Pages. For example, a company – Company A – can protect its LinkedIn Page and only allow job posters at approved email domains (e.g., @companya.com) to post jobs to Company A’s LinkedIn Page. Similarly, a company can protect its Page and only allow job posters from enterprise Recruiter seats post jobs to its

LinkedIn Page. These protections limit the ability of a bad actor to post a fraudulent job that poses as a job from a real company.

- LinkedIn has repeated offender systems to limit the ability of bad actors to repeatedly post jobs that violate LinkedIn’s policies (including fraudulent jobs). LinkedIn has implemented systems at both the job poster and the LinkedIn Page level. That is, a job poster that posts multiple jobs (or in some cases a single job, depending on the policy violation) that violates LinkedIn’s policies will have its job posting ability restricted. Similarly, a LinkedIn Page with job posts that repeatedly violate LinkedIn’s policies will have its job posting ability restricted. This helps limit bad actors from attempting to conduct fraudulent schemes at scale.
- LinkedIn has implemented systems to help educate job seekers regarding how to identify and protect themselves from job scams. In many cases, bad actors execute scams via other systems (e.g., via email, via other messaging platforms).

Human Dignity

Risks related to Human Dignity include exposure to content pertaining to sexual exploitation, pornography, human trafficking, vulgarity, and gore. Absent sufficient mitigations, risks related to Human Dignity may manifest on the platform in ways such as members attempting to share content in Feed or messages that is violent, vulgar, or sexually explicit, or job posters attempting to post a job that facilitates or perpetrates human trafficking.

The **Inherent Probability** of systemic risks concerning Human Dignity stemming from the function, use, or misuse of LinkedIn’s services is assessed as “Likely.” While the overall market demand for mature content and the presence of violence and gore in popular media in Europe is high, there is a lower likelihood of members sharing such content on LinkedIn due to its professional nature. The LinkedIn platform also does not allow for rapid and wide dissemination of content that has not been selected for quality, which limits systemic spread of harmful content. Content shared by members is typically only available to members within their network, which, for any inauthentic accounts that may slip past LinkedIn’s proactive defenses, is likely to be small. Based on the relative frequency of content or activities related to Human Dignity appearing in LinkedIn’s content moderation systems as well as public discourse related to LinkedIn, the probability is assessed as “Likely” that users could share or access information or perpetrate activities negatively impacting Human Dignity on LinkedIn’s platform absent sufficient mitigations.

The **Inherent Severity** of impact for content and activities that negatively impact Human Dignity is rated as “Critical” primarily due to the gravity of risks within this category, considering the potential for significant harm to wellbeing, societal, economic, and security systems at the individual and potentially broader level. Some risks within this category, such as human trafficking and prostitution, are potentially irremediable. Other

risks, such as gore, vulgarity, profanity or pornographic content, may be remediable but can be significant.

With the rating for Inherent Probability as “Likely” and Inherent Severity as “Critical,” the **Inherent Risk** associated with negative impacts to Human Dignity across LinkedIn products and features is therefore assessed as “High.”

In addition to the DTSP Best Practices LinkedIn has effectively implemented across all Risk Areas, which are further described under [“Mitigations,”](#) LinkedIn has implemented mitigations specific to Human Dignity following industry best practices with defined and documented processes. As such, the **maturity of mitigation** efforts LinkedIn has applied to the Human Dignity risk area is assessed as “Defined,” which brings the **Residual Risk** rating down to “Moderate.”

Recognizing that the Residual Risk is not “Low,” LinkedIn continues to prioritize the development and enhancement of strategies to further manage and mitigate risks related to Human Dignity, aiming for an even safer user experience. The key mitigations currently implemented specific to Human Dignity are described in the subsequent paragraphs of this section.

The Professional Community Policies clearly specify what members can and cannot do on LinkedIn. LinkedIn’s Ads and Jobs Policies provide heightened requirements for advertising and job posts on LinkedIn. LinkedIn encourages civil conversation about all topics, including those touching on or relevant to Human Dignity; however, LinkedIn does not allow that conversation to be expressed in an abusive manner. LinkedIn prohibits engagement in sexual innuendos or unwanted advances and the sharing of harmful or shocking material. Content that facilitates prostitution and escort services is not allowed on the platform. LinkedIn’s Ads policies prohibit ads from being offensive to good taste. This means ads must not be, for example, hateful, vulgar, sexually suggestive, or violent. All ads on LinkedIn must comply with LinkedIn’s Professional Community Policies. LinkedIn also does not allow ads that are inappropriate during the time of or in relation to a tragedy, disaster, or similarly sensitive event. LinkedIn removes content that violates the policies and only makes content that meets a higher standard of quality and fits the professional focus of the platform eligible for broader distribution. LinkedIn removes depictions of extreme violence, along with image or video content that is excessively gory, gruesome, or disturbingly shocking. LinkedIn will also warn members when photos or videos that may contain sensitive or graphic content are detected. The content will initially be blurred with a warning, but members can decide if they would like to view it.

If LinkedIn becomes aware that potentially policy-violating content has been posted, LinkedIn reviews that content for policy compliance. Potentially policy-violating content is identified for human review by internally developed classifiers with high levels of precision and recall that are trained on a set of confirmed policy-violating content. In

addition to this proactive review, members can flag potentially policy-violating content for human review by LinkedIn's Trust Review OPS (TRO) team. Content review is conducted by a global team of language and policy enforcement experts who adhere to strict limits with respect to the time in which the review must be completed. In addition, this team is subject to exacting training and quality assurance standards.

LinkedIn's cross-functional Trust team, comprising engineers, product, data science, AI specialists, investigators, and legal teams, work together not only to detect and remove abusive content from LinkedIn, but also to educate LinkedIn members that unwanted sexual advances and other sexual harassment do not belong on LinkedIn. For example, LinkedIn provides guidance through Help Center articles with clear information and examples about the prohibition of sexual innuendos and unwanted advances on the platform, as well as guidance on how to report harassment or safety concerns. LinkedIn also warns members when photos or videos in Feed or Profiles are detected that may contain sensitive or graphic content. The content will initially be blurred with a warning, giving members control over whether they view this content, which can serve a newsworthy or educational purpose while complying with LinkedIn's policies.

LinkedIn offers settings that allow members to only receive communications from people they know. It also offers an optional advanced safety feature in which European members can allow LinkedIn's automated machine learning models to detect sexual harassment and certain other abusive content in messages received by the member, giving the recipient the ability to view and report the message to LinkedIn for review if desired. Where the setting is enabled and LinkedIn's automated systems detect potentially harmful content, depending on whether the sender and recipient have exchanged prior communications, the message will be directed to the recipient's spam folder or be hidden from the recipient by a warning. The warning can be dismissed, giving the recipient the ability to view and report the message to LinkedIn for review if desired. In certain circumstances LinkedIn also may surface a warning and fraud prevention resources to members if its technology detects an attempt to move private conversations away from LinkedIn. This warning is designed to encourage members to consider the likely authenticity of the people with whom they are engaging.

Risk Areas of Low Residual Risk

Considering the probability, severity, and maturity of mitigations, LinkedIn assessed that eight of the assessed risk areas fall into a Residual Risk category of “Low:” Mental and Physical Well-being, Discrimination and Hate, Freedom of Expression and Information, Protection of Personal Data, Civic Discourse and Electoral Processes, Public Health, Illegal Content and Activities, and Private and Family Life.

Mental and Physical Wellbeing

Risks to Mental and Physical Wellbeing include gender-based hate speech and violence, encouraging or assisting suicide, sexual harassment, stalking, encouraging coercive behavior, and addictive behaviors that may impact mental health. Absent sufficient mitigations, risks related to Mental and Physical Wellbeing may manifest on the platform in ways such as members attempting to share content in Feed or Messages that includes threats of violence, topics related to gender-based violence, or language encouraging or assisting suicide, or job postings that lead to harm to user physical or mental wellbeing.

The **Inherent Probability** of systemic risks concerning Mental and Physical Wellbeing stemming from the function, use, or misuse of LinkedIn is assessed as “Likely.” There is a lower likelihood of members sharing content publicly on LinkedIn that would negatively impact Mental and Physical Wellbeing, considering the professional nature of the platform. Additionally, the LinkedIn platform favors knowledge sharing rather than virality, and therefore, tends to not facilitate the rapid and wide dissemination of content that does not meet a minimum quality bar, nor does it contain features that would promote addiction in members within its games or videos. However, considering the relative frequency of content or activities that may negatively impact Mental and Physical Wellbeing appearing in LinkedIn’s content moderation systems as well as public discourse related to LinkedIn, the probability is assessed as “Likely.”

The **Inherent Severity** for content and activities that negatively impact Mental and Physical Wellbeing is rated as “Critical” primarily due to the gravity of risks within this category, considering the potential for significant harm to wellbeing, economic, and societal systems at the individual and potentially broader level. The impact of some risks within this category, such as violence against women, are irremediable. Other risks, such as behavioral addictions, may be considered remediable but still have the potential to cause significant damage to an individual’s mental health or threaten their physical wellbeing.

With the rating for Inherent Probability as “Likely” and Inherent Severity assessed as “Critical,” the **Inherent Risk** associated with negative impacts to Mental and Physical Wellbeing across LinkedIn products and features is therefore assessed as “High.”

In addition to the DTSP Best Practices LinkedIn has effectively implemented across all Risk Areas, which are further described under [“Mitigations,”](#) LinkedIn has implemented

robust mitigations following industry best practices with defined, documented, and managed processes. As such, the **maturity of mitigation** efforts LinkedIn has applied to the Mental and Physical Wellbeing risk area is assessed as “Managed,” which brings the **Residual Risk** rating down to “Low.”

Considering the High Inherent Risk related to Mental and Physical Wellbeing, LinkedIn continues to prioritize the development and enhancement of strategies to further manage and mitigate risks related to Mental and Physical Wellbeing, aiming to maintain a high level of user safety on the platform. The key mitigations currently implemented specific to Mental and Physical Wellbeing are described in the subsequent paragraphs of this section.

As a professional network, LinkedIn maintains a high standard for safe and civil conversations between members. LinkedIn policies prohibit a wide range of actions and behaviors associated with harassment, bullying, or other abusive conduct. LinkedIn prohibits gender-based hate speech, symbols, and groups on the platform, and removes any content that attacks, denigrates, intimidates, dehumanizes, incites or threatens hatred, violence, or prejudicial or discriminatory action against individuals or groups because of gender, gender identity, or sexual orientation. LinkedIn also removes content that personally attacks, intimidates, shames, bullies, disparages other LinkedIn members, or is otherwise abusive or uncivil. LinkedIn does not allow the sharing of harmful or shocking material. LinkedIn may permit content that evokes hateful rhetoric (including slurs) in the context of counter speech, reclamation, or members’ personal experiences with sexism. LinkedIn also allows members to call out actual or perceived hateful, prejudicial, or discriminatory conduct, provided such content doesn’t otherwise violate its policies. In connection with a newsworthy event, LinkedIn may label such content instead of removing it from the platform. LinkedIn will not penalize members for posting content in these circumstances but may take steps to limit its distribution to protect other members who may not wish to see it.

If a member is sharing dismissive or derisive content within the LinkedIn Feed - content that discourages, mocks, ridicules, makes fun of or refutes other individuals or groups or their self-identification or identity - this content may not be eligible for broad distribution on the platform. Additionally, if other members are expressing disinterest in, dismissing or reporting a certain piece of content, that content may not receive broad distribution on the platform, as this indicates members potentially find the content to be encouraging gender-based violence or otherwise unhelpful or unwanted.

If LinkedIn becomes aware that potentially policy-violating content has been posted, LinkedIn reviews that content for policy compliance. Potentially policy-violating content is identified for human review by internally developed classifiers with high levels of precision and recall that are trained on a set of confirmed policy-violating content. In addition to this proactive review, members can flag potentially policy-violating content

for human review by LinkedIn's Trust Review OPS (TRO) team. Content review is conducted by a global team of language and policy enforcement experts who adhere to strict limits with respect to the time in which the review must be completed. In addition, this team is subject to exacting training and quality assurance standards.

LinkedIn offers settings that allow members to only receive communications from people they know. It also offers an optional advanced safety feature in which members can allow LinkedIn's automated machine learning models to detect sexual harassment and certain other abusive content in messages received by the member, giving the recipient the ability to view and report the message to LinkedIn for review if desired. Where the setting is enabled and LinkedIn's automated systems detect potentially harmful content, depending on whether the sender and recipient have exchanged prior communications, the message will be directed to the recipient's spam folder or be hidden from the recipient by a warning. The warning can be dismissed, giving the recipient the ability to view and report the message to LinkedIn for review if desired.

LinkedIn also warns members when photos or videos in Feed or Profiles are detected that may contain sensitive or graphic content. The content will initially be blurred with a warning, giving members control over whether they view this content, which can serve a newsworthy or educational purpose while complying with LinkedIn's policies.

Members can manage the visibility of their information through settings and also who their data is shared with (i.e., post visibility). These settings are honored throughout the LinkedIn platform, including on Feed, Search, and other personalized modules like People You May Know and Notifications. For example, members can:

- Control how their Profile is shown on LinkedIn, such as on search engines.
- Control who can send them an invitation to connect on LinkedIn (e.g., members can choose to receive connection invitations only from people who already know their email address);
- Disallow non-connections from sending them messages (to avoid unwanted contact);
- Block particular members by using the "block" feature (see more here);
- Restrict the visibility of their email address, connections list, and last name; and
- Disallow others from being able to find their LinkedIn Profile using the member's email address or phone number.

Discrimination and Hate

Risks related to Discrimination and Hate include exposure to discriminatory content or activities on the platform, including hate speech, and potential bias and discrimination in LinkedIn's algorithmic systems. Absent sufficient mitigations, risks related to Discrimination and Hate may manifest on the platform in ways such as members attempting to share content in Feed that contains hate speech, job posters attempting to post a job that discriminates against certain races or genders, or LinkedIn's recommender systems recommending candidates to recruiters based on bias in algorithms.

The **Inherent Probability** of systemic risks concerning Discrimination and Hate stemming from the function, use, or misuse of LinkedIn's services is assessed as "Likely." There is a lower likelihood of members sharing content publicly on LinkedIn related to Discrimination and Hate, considering the professional nature of the platform. Additionally, the nature of the LinkedIn platform favors knowledge sharing rather than virality, and therefore, tends to not facilitate the rapid and wide dissemination of content that does not meet a minimum quality bar. In turn, this limits the potential for discriminatory or hateful content to spread in a viral nature on the platform. Absent sufficient mitigations, LinkedIn's recommender systems could operate with bias or make recommendations with discriminatory implications. Considering the relative frequency of content or activities related to Discrimination and Hate appearing in LinkedIn's content moderation systems as well as public discourse related to LinkedIn, the probability is assessed as "Likely."

The **Inherent Severity** for content and activities that negatively impact Discrimination and Hate on LinkedIn is rated as "High" primarily due to the gravity of risks within this category, considering the impact to wellbeing, societal, and economic impacts at the individual level and societal impacts up to a regional scale. Risks within this category, such as hate speech and discrimination, may be remediable but can cause serious damage, including loss of career advancement opportunities.

With the rating for Inherent Probability as "Likely" and Inherent Severity as "High," the **Inherent Risk** associated with Discrimination and Hate across LinkedIn products and features is therefore assessed as "High."

In addition to the DTSP Best Practices LinkedIn has effectively implemented across all Risk Areas, which are further described under "[Mitigations](#)," LinkedIn has implemented robust mitigations following industry best practices with defined, documented, and managed processes. As such, the **maturity of mitigation** efforts LinkedIn has applied to the Discrimination and Hate risk area is assessed as "Managed," which brings the **Residual Risk** rating down to "Low."

LinkedIn has invested resources into ensuring "two members who are equally qualified should have equal access to opportunity" on LinkedIn, as well as limiting hate speech

on the platform. LinkedIn continues to prioritize the development and enhancement of strategies to further manage and mitigate risks related to Discrimination and Hate on the platform. The key mitigations currently implemented specific to Discrimination and Hate are described in the subsequent paragraphs of this section.

LinkedIn policy prohibits hate speech, symbols, and groups on the platform and removes any content that attacks, denigrates, intimidates, dehumanizes, and incites or threatens hatred, violence, prejudicial or discriminatory action against individuals or groups because of their actual or perceived race, ethnicity, national origin, caste, gender, gender identity, sexual orientation, religious affiliation, or disability status. LinkedIn also removes content that personally attacks, intimidates, shames, bullies, or disparages other LinkedIn members, or is otherwise abusive or uncivil. Additionally, members may not engage in trolling or other repetitive negative content that disrupts other members or conversations on the platform.

LinkedIn's contracts and policies prohibit discriminatory behavior. LinkedIn's policies are set forth in its LinkedIn Jobs Terms and Conditions and LinkedIn Service Terms.

LinkedIn developed classifiers with high levels of precision and recall that are trained on a set of confirmed policy-violating content identified by human reviewers. Such classifiers are used to proactively enqueue potentially policy-violating content for human review. In addition to proactive review, it is straightforward for members to flag potentially policy-violating content for human review. Content review is conducted by a global team of language and policy enforcement experts that adhere to strict limits with respect to the time in which the review must be completed. In addition, this team is subject to exacting training and quality assurance standards.

If a member is sharing dismissive or derisive content within the LinkedIn Feed - content that discourages, mocks, ridicules, makes fun of or refutes other individuals or groups or their self-identification or identity - this content may not be eligible for broad distribution on the platform. Additionally, if other members are expressing disinterest in, dismissing or reporting a certain piece of content, that content may not receive broad distribution on the platform, as this indicates members potentially find the content to be discriminatory or otherwise unhelpful or unwanted.

For Jobs, LinkedIn prohibits discrimination in job posts based on protected characteristics, including age, gender, gender identity, religion, ethnicity, race, national origin, disability, sexual orientation, and any other basis protected under law; however, in a jurisdiction where it is legally permissible to do so, LinkedIn may permit language in posts expressing preference for members of certain groups historically disadvantaged in hiring in that location. LinkedIn automatically reviews job postings created by members for potential discrimination and flags potentially violative jobs for human review. If a decision is made that a specific job posting is discriminatory, LinkedIn searches for and

removes other copies of the job posting on the system. Members can also report for review by LinkedIn any job post on LinkedIn that they believe may be discriminatory.

In-product functionality is also utilized within Jobs to deter member posting of discriminatory content, such as an inclusivity nudge that brings potentially discriminatory language in draft job posts to posters' attention before submission, so they can avoid posting jobs that run afoul of LinkedIn's Jobs Policies. Filters on Recruiter do not allow customers to search for members along with any protected characteristics. Customers can also enable in-product functionality like Hide Candidate Photos and Names. Additionally, all jobs can only be targeted (i.e., promoted to LinkedIn members) by LinkedIn customers based on very limited criteria – general region and skills. Moreover, all jobs are visible to all members, regardless of targeting criteria identified by the job poster for promotion. In other words, all members could still find a particular job post on LinkedIn even if it were not specifically promoted to them (because of skills or location).

If an advertiser attempts to target an ad based on age or gender, LinkedIn first requires the advertiser to certify that they will not use LinkedIn to discriminate based on age, gender, or other protected characteristics in ads relating to employment, housing, education, and credit, and provides a link to a Help Center article detailing LinkedIn's policy. LinkedIn also has specific policies preventing ad targeting discrimination, as well as policies applicable to particular types of ads products on LinkedIn. For example, terms of use for Lead Gen Forms prohibit the collection of sensitive data. LinkedIn also prohibits the use of Insight Tag on Pages collecting specific health-related or financial related webpages (to protect against the inadvertent collection of data that could be used to facilitate discrimination). LinkedIn does not allow targeting on protected characteristics other than age and gender and only allows ad targeting based on age or gender with the advertiser first providing this certification. To protect against abuse through exclusion, customers are not able to exclude ads based on age or gender.

LinkedIn employs in-product functionality designed to prevent and deter potentially discriminatory behavior on the part of our customers and systems. "Lookalike" audience features (e.g. predictive or expanded audiences) do not use sensitive demographic attributes, including attributes such as sex or age within the model.

As detailed earlier, Ads cannot be targeted based on any of the protected categories except age / gender. LinkedIn does not allow ad targeting of members in Europe using LinkedIn Groups. If a customer opts to utilize demographic targeting like age or gender, within Campaign Manager, LinkedIn advertisers are also required to certify via a checkbox that if an ad relates to employment, housing, education or credit, they will not use LinkedIn to discriminate based on age or other protected characteristics. Only once an advertiser checks this box will the option to target by age or gender become available. If the stated objective for a campaign is talent leads and recruitment, the option to target by age is not available at all.

Before using Insight Tag, customers must agree that they "will not install the Insight Tag ... on web pages that collect or contain sensitive data," including "pages offering specific health-related or financial services to consumers."

LinkedIn offers settings that allow members to only receive communications from people they know. It also offers an optional advanced safety feature in which European members can allow LinkedIn's automated machine learning models to detect sexual harassment and certain other abusive content in messages received by the member, giving the recipient the ability to view and report the message to LinkedIn for review if desired. Where the setting is enabled and LinkedIn's automated systems detect potentially harmful content, depending on whether the sender and recipient have exchanged prior communications, the message will be directed to the recipient's spam folder or be hidden from the recipient by a warning. The warning can be dismissed, giving the recipient the ability to view and report the message to LinkedIn for review if desired. In certain circumstances LinkedIn also may surface a warning and fraud prevention resources to members if its technology detects an attempt to move private conversations away from LinkedIn. This warning is designed to encourage members to consider the likely authenticity of the people with whom they are engaging.

LinkedIn is committed to measuring and preventing bias and discrimination in its algorithmic systems. Recommender systems including the LinkedIn Feed are assessed for unintended bias. Currently, that is only done across inferred binary gender. LinkedIn's approach to Responsible AI is forward leaning. LinkedIn examines the ethics and compliance of its AI-powered products and services as part of the Trustworthy Design Review (TDR) process, which it applies to LinkedIn product initiatives prior to launch. LinkedIn has also developed tools to enable the measurement of fairness in large-scale machine learning and artificial intelligence workflows. They can be deployed in training and scoring workflows to measure biases in training data, evaluate different fairness notions for ML models, and detect statistically significant differences in model performance across different subgroups. LinkedIn conducts fairness checks on Feed distribution and ranking. LinkedIn evaluates its algorithms for binary gender bias, relying on a combination of member provided gender data (through [Self-ID](#)) and binary gender inference. Additionally, LinkedIn Talent Solutions systems including Recruiter search, Recommended Matches, and Jobs You Might Be Interested In (JMBII) are regularly assessed for unintended bias.

LinkedIn also conducts Fairness Checks on Job recommendations to ensure jobs are shown equally across demographics based on available data. LinkedIn conducts A/B testing of algorithmic systems and leverages long term holdouts for portions of the LinkedIn population to identify any potential undesirable impacts resulting from changes to algorithmic systems. LinkedIn also monitors for unexpected anomalies in distribution of content to track and avoid the introduction of bias into LinkedIn's algorithms. To prevent job discrimination, LinkedIn launched a Discrimination Prevention Model across

all job types (including online, field, and ingested jobs) to improve model precision and recall. Now more discriminatory jobs will be detected and prevented, resulting in reduced harm to seekers (abusive views on discriminatory jobs) and more equitable hiring marketplace outcomes. LinkedIn employs repeat offender systems so that members that repeatedly post policy-violating content face escalating consequences, including potential permanent restriction from LinkedIn.

Freedom of Expression and Information

Risks to Freedom of Expression and Information include risks that freedom of expression, information, the press, or pluralism of the media are negatively impacted by the use, design, or functioning of a platform or its algorithmic systems, policies, or enforcement. Absent sufficient mitigations, risks related to Freedom of Expression and Information may manifest on the platform in ways such as LinkedIn's content moderation systems removing content that constitutes legally permissible speech, preventing legitimate jobs from being posted, or amplifying content from some but not all media viewpoints.

The **Inherent Probability** of systemic risks concerning Freedom of Expression and Information stemming from the use, misuse, or function of LinkedIn's services is assessed as "Likely." As a real-identity, specific-purpose platform, LinkedIn operates under standards of professionalism, which are reflected both in content policies and enforcement, as well as in content prioritization and amplification. LinkedIn strictly enforces both policies for removing content from the platform and for determining when content is eligible for broader distribution. Therefore, absent mitigations to ensure fair and appropriate application of these content moderation policies, the probability of negative impacts to user rights to Freedom of Expression and Information is assessed as "Likely."

The **Inherent Severity** for content and activities that negatively impact Freedom of Expression and Information on LinkedIn is rated as "High" primarily due to the gravity of risks within this category, considering the potential for significant harm to societal and wellbeing systems up to the country level and political systems up to the regional level. Depending on the local political environment, consequences related to the negative impact on Freedom of Expression and Information can be irremediable. But the impact of many risks within this category is not considered significant.

With the rating for Inherent Probability as "Likely" and Inherent Severity as "High," the **Inherent Risk** associated with negative impacts to Freedom of Expression and Information across LinkedIn products and features is therefore assessed as "High."

In addition to the DTSP Best Practices LinkedIn has effectively implemented across all Risk Areas, which are further described under "[Mitigations](#)," LinkedIn has implemented a robust set of mitigations following industry best practices with defined, documented, and managed processes. As such, the **maturity of mitigation** efforts LinkedIn has

applied to the Freedom of Expression and Information risk area is assessed as “Managed,” which brings the **Residual Risk rating** for this area down to “Low.”

LinkedIn carefully weighs trade-offs between freedom of expression and online harms, erring on the side of freedom of expression where possible. Over the review period, authors appealed only about 4.26% of the content moderation decisions that LinkedIn made. LinkedIn granted roughly 22.63% of those appeals, representing 1.49% of moderation decisions. These numbers illustrate the priority LinkedIn places on freedom of expression considering 1) LinkedIn’s receptivity to a meaningful percentage (22.63%) of appeal requests and 2) the very small percentage of total content moderation decisions overturned (1.49%). That said, LinkedIn continues to invest in developing and enhancing strategies to further mitigate risks to Freedom of Expression and Information. The key mitigations currently implemented specific to Freedom of Expression and Information are described in the subsequent paragraphs of this section.

If LinkedIn becomes aware that potentially policy-violating content has been posted, LinkedIn reviews that content for policy compliance. Potentially policy-violating content is identified for human review by internally developed classifiers with high levels of precision and recall that are trained on a set of confirmed policy-violating content. In addition to this proactive review, members can flag potentially policy-violating content for human review by LinkedIn’s Trust Review OPS (TRO) team.

LinkedIn creates value and preserves trust by shielding members from harmful content and abuse, while honoring their professional expression and speech. LinkedIn enables healthy on-platform conversations by facilitating the removal of content that threatens members’ safety and creates barriers to their engagement on LinkedIn. When speech does not conclusively cross the boundaries set by LinkedIn policies, LinkedIn minimizes interference and gives the speaker the benefit of the doubt and, subject to certain limited exceptions, generally will not remove or geo-block the visibility of the content without valid legal process requiring it to do so.

Content review is conducted by a global team of language and policy enforcement experts who adhere to strict limits with respect to the time in which the review must be completed. In addition, this team is subject to exacting training and quality assurance standards. LinkedIn has clear and well-documented policies and tracks edge cases for sensitive escalations to enhance or clarify policies as needed. LinkedIn employs a dedicated team of trainers, who not only support the onboarding of new content moderators, but also provide ongoing educational opportunities for all moderators. Moderators have direct access to global content policy managers through regular office hours and dedicated escalation pathways. For particularly complex decisions, content policy managers also have access to in-house lawyers who can consult regional legal experts as needed.

Users have access to report content on each LinkedIn surface. Both reporter and creator generally are notified when content is removed, and both are given an opportunity to appeal the decision.

LinkedIn responds to periodic evaluations by the third-party independent organization Ranking Digital Rights before it publishes its annual ratings on LinkedIn practices, governance, and leadership on the protection of freedom of expression and privacy. LinkedIn also engages with its parent company, Microsoft, in connection with Microsoft's commitments to human rights, including free expression.

LinkedIn carefully considers all government removal requests and works to mitigate any implications they may have on freedom of expression and human rights. For government demands, LinkedIn employs safeguards to ensure any actions taken are narrow, specific, submitted in writing, and based on valid legal orders. Through its parent company, Microsoft, LinkedIn also engages with broader civil society organizations on best practices related to government removal requests and participates in human rights impact assessments. LinkedIn generally provides prompt notice to the member(s) whose content is at issue in a government removal request, unless legally prohibited from doing so. This notice may enable the impacted member to challenge the request served on LinkedIn.

LinkedIn is also committed to measuring and preventing bias in its algorithmic systems that would negatively impact freedom of expression and pluralism. LinkedIn has developed tools to enable the measurement of fairness in large-scale machine learning and artificial intelligence workflows. They can be deployed in training and scoring workflows to measure biases in training data, evaluate different fairness notions for ML models, and detect statistically significant differences in model performance across different subgroups. LinkedIn conducts fairness checks on Feed distribution and ranking. LinkedIn evaluates its algorithms for binary gender bias, relying on a combination of member provided gender data (through Self-ID) and binary gender inference. LinkedIn also monitors for unexpected anomalies in distribution of content to track and ensure the introduction of bias into LinkedIn's algorithms.

Protection of Personal Data

Risks to Protection of Personal Data include the collection, processing, mishandling, or disclosure of user data without a lawful basis, including data mining, profile scraping, phishing, hacking, malware, and data breaches. Absent sufficient mitigations, risks related to Protection of Personal data may manifest on the platform in ways such as LinkedIn inappropriately using personal data without a lawful basis, inadequate transparency in how LinkedIn processes and stores user data, or improper leakage of data to the public.

The **Inherent Probability** of systemic risks concerning Protection of Personal Data is assessed as “Likely.” Although a high level of sophistication is required to perpetrate harm such as hacking, malware, or data breaches, as well as internal misuse of member data, the financial incentives of fraudulent activities lead to a higher demand for theft of personal data. Considering the relative frequency of content or activities related to Human Dignity appearing in LinkedIn’s content moderation systems as well as public discourse related to LinkedIn, the probability is assessed as “Likely.”

The **Inherent Severity** for content and activities that negatively impact Protection of Personal Data is rated as “High” primarily due to the gravity of risks within this category, considering the potential for significant harm to economic, security, and societal systems at the individual and local level. Some risks within this category – particularly related to phishing, hacking, malware, or data breaches – may reach a broader scale of impact while others, such as collection of data without consent, may only impact the individual. For the most part, these risks are remediable but may include some irremediable damage.

With the rating for Inherent Probability as “Likely” and Inherent Severity as “High,” the **Inherent Risk** associated with Protection of Personal Data across LinkedIn products and features is therefore assessed as “High.”

In addition to the DTSP Best Practices LinkedIn has effectively implemented across all Risk Areas, which are further described under [“Mitigations,”](#) LinkedIn has implemented robust mitigations specific to the Protection of Personal Data following industry best practices with defined, documented, and managed processes. As such, the **maturity of mitigation** efforts LinkedIn has applied to the Protection of Personal Data risk area is assessed as “Managed,” which brings the **Residual Risk rating** down to “Low.”

LinkedIn takes the responsibility of protecting user data very seriously and is actively working to develop and enhance strategies to further manage and mitigate risks related to Protection of Personal Data, aiming for even greater protection for users. LinkedIn continues to invest in policies, policy enforcement mechanisms, and advanced detection capabilities, including capabilities to prevent Profile scraping and address other evolving threats related to user privacy. The key mitigations currently implemented specific to Protection of Personal Data are described in the subsequent paragraphs of this section.

LinkedIn’s core value is “members first.” That value guides LinkedIn as it aims to provide clarity to platform members, act consistently, and give its members control. That means LinkedIn aims to be transparent in its policies and offers its members settings to help them control the data they provide on the platform. LinkedIn is compliant with the General Data Protection Regulation (GDPR) and allows users worldwide to access, correct errors in, erase, object to the processing of, and export their personal data. LinkedIn also requires its customers, suppliers, data processors, and sub processors to enter into Data Protection Agreements (“DPAs”) that contractually limit how the

customer must handle data. The DPAs require data to be handled in accordance with GDPR and other applicable laws and include EU approved Standard Contractual Clauses which permit the transfer of data outside of the EU. Members also can exercise their data subject rights by accessing, objecting to, and managing LinkedIn's use of their data or they can delete it from LinkedIn. LinkedIn has appointed a Data Protection Officer (DPO) to ensure that its processing of personal data is in compliance with GDPR. The DPO is the 'voice of the member' within LinkedIn. He makes recommendations in relation to new and existing data processing operations and handles complaints relating to the processing of personal data.

LinkedIn's Privacy Policy and Help Center Articles provide members with notice about the data that LinkedIn collects and how LinkedIn handles their personal data. The Privacy Policy provides an overview, while the Help Center Articles provide a more in-depth explanation for specific features. LinkedIn uses a variety of in-product indicators to enhance member awareness and empowerment, including "Why am I seeing this ad?" (WAISTA).

Trust is an operating priority at LinkedIn that includes Safety, Privacy, and Security and is reported on at regular company-wide meetings and directly to the Chief Executive Officer. LinkedIn has invested in privacy protections and privacy enhancing technologies through the establishment of a Data Privacy Engineering team that works in line with the Privacy Legal team. A Privacy Impact Assessment is built into the review process for initiatives or features assessed as potentially likely to include personal data.

LinkedIn has robust practices to safeguard and cultivate a culture of respect for user privacy. LinkedIn has its own internal Data Privacy Principles for data processing across the business and maintains an InfoSec standard which supports the proper handling of data within LinkedIn's data ecosystem. LinkedIn maintains a robust global information security team that protects its infrastructure and member data contained therein. This team is comprised of world-class technologists and investigators that ensure that anomalies (identified through automation) are investigated to ensure that external actors cannot access LinkedIn's internal systems.

For higher-risk activities that involve the processing of personal data, LinkedIn completes Data Protection Impact Assessments ("DPIAs"). DPIAs document the type of data collected (data elements and whether sensitive data is included), risks associated with the processing (including impact to the fundamental rights and freedoms of data subjects), and mitigations that have been implemented to reduce such risks. Once completed, LinkedIn's Data Privacy Officer (DPO) reviews the DPIA, asks independent questions to ensure that the activity is in line with GDPR and member expectations, provides recommendations, and signs off once the risk is deemed to be acceptable.

To boost account security for members and Job posters, LinkedIn encourages the enablement of two-factor authentication (2FA). Clicking "Secure account" will redirect members to the settings page to complete enrollment.

To protect members' privacy, LinkedIn prohibits the use of third-party software or browser extensions that scrape, modify the appearance of, or automate activity on LinkedIn's website.

LinkedIn also employs a robust set of mitigations to identify and address external threats to the Protection of Personal Data, which have been described in further detail in the [Consumer Protection and Fraud section](#) of this Report. Most notably, LinkedIn reviews public-facing content for harmful or abusive indicators and flags potential violative content for review. It scans both public- and privately-facing content for phishing and malware that could offer bad actors access to member data.

Civic Discourse and Electoral Processes

Risks to Civic Discourse and Electoral Processes include election misinformation, disinformation, coordinated manipulation and inauthentic activity, foreign interference, and the disruption of electoral processes, as well as political polarization, political bias, and the creation of echo chambers. Absent sufficient mitigations, risks related to Civic Discourse and Electoral Processes may manifest on the platform in ways such as members posting inaccurate or misleading information on electoral processes to Feed, advertisements hyper-targeting vulnerable groups with election misinformation, or recommender systems recommending political content in a biased manner or a manner that contributes to political polarization or the creation of echo chambers.

The **Inherent Probability** of systemic risks concerning Civic Discourse and Electoral Processes stemming from the function of LinkedIn's services is assessed as "Not Likely." There is a lower likelihood of members sharing content that could negatively impact Civic Discourse and Electoral Processes on LinkedIn due to its professional nature. The LinkedIn platform favors knowledge sharing rather than virality, and therefore, tends to not facilitate the rapid and wide dissemination of content that does not meet a minimum quality bar. In turn, this limits the systemic risk posed by the spread of harmful content. LinkedIn also does not allow political advertisements on the platform. Considering these factors and the relative frequency of content or activities related to Civic Discourse and Electoral Processes appearing in LinkedIn's content moderation systems as well as public discourse related to LinkedIn, the probability is assessed as "Not Likely" that users could share or be exposed to information negatively impacting Civic Discourse and Electoral Processes on the platform.

The **Inherent Severity** for content and activities that negatively impact Civic Discourse and Electoral Processes on LinkedIn is rated as "Critical" primarily due to the gravity of risks within this category, considering the potential for significant harm to political, societal, economic, and security systems at the local, country, and even regional levels.

Some risks within this category have potentially irremediable consequences on political or security systems. The impact of other risks may be considered remediable but still have the potential to cause significant damage.

With the rating for Inherent Probability as “Not Likely” and Inherent Severity as “Critical,” the **Inherent Risk** associated with negative impacts to Civic Discourse and Electoral Processes across LinkedIn products and features is therefore assessed as “Moderate.”

In addition to the DTSP Best Practices LinkedIn has effectively implemented across all Risk Areas, which are further described under “[Mitigations](#),” LinkedIn has implemented robust mitigations specific to Civic Discourse and Electoral Processes, following industry best practices with defined, documented, and managed processes. LinkedIn also considered the guidance provided by the European Commission (Commission Guidelines for providers of Very Large Online Platforms and Very Large Online Search Engines on the mitigation of systemic risks for electoral processes pursuant to Article 35(3) of Regulation (EU) 2022/2065) in evaluating the maturity of mitigations for this Risk Area. While many of LinkedIn’s existing practices aligned with the guidelines, LinkedIn made additional enhancements during this review period to risk mitigations for election integrity consistent with the guidelines. As such, the **maturity of mitigation** efforts LinkedIn has applied to the Civic Discourse and Electoral Processes risk area is assessed as “Managed,” which brings the **Residual Risk** rating down to “Low.”

The key mitigations currently implemented to address risks to Civic Discourse and Electoral Processes are described in the subsequent paragraphs of this section.

LinkedIn encourages civil conversation about all topics on the platform, including topics touching on Civic Discourse and Electoral Processes. However, LinkedIn does not allow that conversation to be expressed in a hateful or bullying manner. LinkedIn does not allow users to share false or misleading content or content to interfere with or improperly influence an election or other civic process. LinkedIn also prohibits political advertisements on its platform, including ads advocating for or against a candidate, party, or ballot proposition, or ads designed for political fundraising or otherwise intended to influence an election outcome.

If potentially policy-violating content is posted, LinkedIn reviews that content for policy compliance. Potentially policy-violating content is identified for human review through classifiers trained through machine learning and member reports. Internally developed classifiers proactively enqueue potentially policy violating content for human review. In addition to proactive review, it is straightforward for members to flag potentially policy-violating content for human review. Content review is conducted by a global team of language and policy enforcement experts that adhere to strict limits with respect to the time in which the review must be completed. The review team is subject to exacting training and quality assurance standards.

LinkedIn aims to educate its members about civic discourse, electoral processes, and public security through its global team of news editors. LinkedIn has provided a [Help Center article](#) to support media literacy, and LinkedIn has an internal team of global news editors that provides trustworthy and authoritative content to its member-base at all times.

During the pre-election cycle, LinkedIn relies on trusted and reputable publisher sources for featured shares, focusing on the policy impact on businesses and professionals around the EU. LinkedIn also curates links to topical landing pages from trusted publishers to provide members with easy and reliable entry points to more detailed coverage. LinkedIn does not compete with trusted publishers for speed or depth of coverage but instead aims to connect their existing coverage to LinkedIn members and their needs. During the European Elections, this team provided manually curated and localized storylines.

LinkedIn has a Nation State Threat program to account for state-sponsored malign actor relationships to content on the platform. LinkedIn's team of threat investigators and intelligence analysts works with peers and other stakeholders, including the AI modeling team, to identify and remove nation-state actors and coordinated inauthentic campaigns. LinkedIn conducts investigations into election-related influence operations and nation-state targeting including continued information sharing on threats with industry peers and Law Enforcement on a regular basis. LinkedIn works with peer companies and other stakeholders to receive and share indicators related to fake accounts created by state-sponsored actors, such as confirmed Tactics, Techniques, and Protocols (TTPs) and Indicators of Compromise (IOC). This exchange of information leads to a better understanding of the incentives of sophisticated and well-funded threat actors and how they evolve their TTPs to achieve those goals, which assists LinkedIn in their identification and removal. LinkedIn's teams also work closely with Microsoft's Threat Intelligence Center (MSTIC) and Democracy Forward teams on security issues.

LinkedIn has a dedicated Civic Integrity team that works to operationalize the additional guidance set forth from the EC around election integrity. The Civic Integrity team is global and cross-functional, with representation across the organization including the Legal, Public Policy, Content Policy, Communications, Escalations, Enforcement Operations, and Product teams, as well as the DSA Compliance Function.

LinkedIn has developed a quantitative and qualitative approach to assessing the priority of global elections to adequately and consistently stand up enhanced internal coverage of election risk based on relative impact. This assessment is used to determine the length and severity of the risk mitigation measures that need to be implemented for each high priority election, including the language resourcing requirements. As a result of this analysis and as needed, LinkedIn may stand up a cross-functional tiger team in

support of certain elections including individuals with country-specific expertise. LinkedIn has also developed incident response plans and escalation paths to address potential crisis events.

LinkedIn continues to invest in its elections metrics framework to more efficiently identify election-related content. LinkedIn monitors conversation via keywords during high priority elections to identify violations or when operational capacity breaches a threshold and requires additional support. LinkedIn established a Multi-Level Review Protocol for election-related content. Public content on LinkedIn is scanned for election-related information and undergoes additional human review as needed to ensure a high level of election integrity. LinkedIn works with external fact checkers and leverages a curated repository of reputable global fact-checking resources across multiple priority languages and localities to support accurate review of election-related content.

LinkedIn removes specific claims, presented as fact, that are demonstrably false or substantially misleading and likely to cause harm. LinkedIn also removes content that is synthetic or manipulated in a way to misrepresent or distort real-life events without clear disclosure of the fake or altered nature of the material. Content that is false or substantially misleading but not likely to cause harm is not eligible for distribution beyond the author's network.

LinkedIn also engages periodically with Member State and EU-level officials to provide information about its election preparedness safety measures, including the steps it takes per the European Commission's election guidelines to mitigate systemic risks for election purposes. Additionally, as part of LinkedIn's engagement with the Rapid Response Group as part of the Code of Practice and in view of the upcoming Presidential Election in Croatia on 29 December 2024, LinkedIn participated in the targeted email communication channel between relevant signatories and members of the task-force to ensure important information and feedback can be efficiently shared regarding potential threats and content posing risks to the electoral processes.

During a portion of the assessment period, LinkedIn had a cross-functional Germany tiger team in place to monitor and evaluate election-related content on LinkedIn which is minimal with only about 100 pieces of content being reviewed daily. Prior to the national election in Germany on 23 February 2025, LinkedIn monitored potentially policy violating content and no major issues or content violations were detected along with minimal escalation activity. In addition, the Trust Product team implemented global mitigations to prevent false positive restrictions caused by sudden spikes in account block volume which can be triggered by mass messages related to elections.

Finally, to respond to the ever-changing threat landscape, LinkedIn's teams continually invest in new technologies for combating inauthentic behavior on the platform. For example, LinkedIn is investing in AI technologies such as advanced network algorithms that detect communities of fake accounts through similarities in their content and

behavior; computer vision and natural language processing algorithms for detecting AI-generated elements in fake Profiles such as deep fakes; anomaly detection of risky behaviors; and deep learning models for detecting sequences of activity that are associated with abusive automation. LinkedIn has also adopted the Coalition for Content Provenance and Authenticity's industry-leading "Content Credentials" technology (C2PA) to include metadata labeling, including data about whether content is created using AI, on content containing the C2PA technology. Furthermore, LinkedIn acts vigilantly to maintain the integrity of all accounts and to ward off false account activity:

- Monitoring platform conversations regarding significant elections and establishing metrics for when election-related conversations, violations, or operational capacity breach a threshold and require additional support;
- Maintaining a dedicated Anti-Abuse team to research emerging trends and key risks and develop tools to address them;
- Using AI to detect inauthentic activity and communities of fake accounts;
- Conducting hash matching for known instances of deepfake content;
- Using automated systems to detect and block automated activity; and
- Maintaining 24/7 escalation paths to address any emerging issues.

Public Health

Risks to Public Health include false information and the promotion of harmful, yet legal, substances or practices (e.g. weight loss medications). Absent sufficient mitigations, risks related to Public Health may manifest on the platform in ways such as members posting misleading health-related information, News results including low quality information related to health events, or LinkedIn Ads promoting legal but harmful substances.

The **Inherent Probability** of systemic risks concerning Public Health stemming from the function, use, or misuse of LinkedIn's services is assessed as "Not Likely." The LinkedIn platform favors knowledge sharing rather than virality, and therefore, tends to not facilitate the rapid and wide dissemination of content that does not meet a minimum quality bar. In turn, this limits the potential for health misinformation to spread in a viral nature on the platform. Considering the relative frequency of content or activities related to Public Health appearing in LinkedIn's content moderation systems as well as public discourse related to LinkedIn, which has decreased this review period, the probability that members may share or be exposed to misleading health information is assessed as "Not Likely."

The **Inherent Severity** for content and activities that negatively impact Public Health on LinkedIn is rated as "Critical," primarily due to the gravity of risks within this category, considering the potential for significant and irremediable harm to wellbeing, economic, and societal systems at the individual and up to global level. Some risks within this

category such as exposure to harmful substances have potentially irreparable consequences.

With the rating for Inherent Probability as “Not Likely” and Inherent Severity as “Critical,” the **Inherent Risk** associated with the negative impacts to Public Health across LinkedIn products and features is therefore assessed as “Moderate.”

In addition to the DTSP Best Practices LinkedIn has effectively implemented across all risk areas, which are further described under “[Mitigations](#),” LinkedIn has implemented mitigations following industry best practices with defined and documented processes. As such, the **maturity of mitigation** efforts LinkedIn has applied to the Public Health risk area is assessed as “Defined,” which brings the **Residual Risk rating** down to “Low.”

LinkedIn continues to prioritize the development and enhancement of strategies to further manage and mitigate risks related to Public Health, aiming for a safer user experience. The key mitigations currently implemented specific to Public Health are described in the subsequent paragraphs of this section.

LinkedIn has invested in new technologies to combat inauthentic behavior that may contribute to risk to Public Health and has robust policies and protocols in place around healthcare advertising and related information.

LinkedIn has policies related to healthcare advertisements that limit drug advertisements to approved retailers. Marketing Solutions’ policies only allow drug advertising in geographic jurisdictions where it is legal. LinkedIn Marketing prohibits the promotion of opioids, anabolic steroids, or other unsafe products on LinkedIn’s prohibited list. LinkedIn leverages a specialized vendor to certify pharmacy and telehealth service ads and limits their distribution to within the U.S. only. LinkedIn also uses classifiers to identify and prohibit content related to opioids, anabolic steroids, or other unsafe products on Feed and other surfaces.

LinkedIn policies prohibit direct-to-purchase links and/or private contact information to facilitate the sale of prescription drugs, pharmaceuticals, alcohol, and tobacco products. LinkedIn has developed classifiers to automatically detect this content or health misinformation on the platform and has also partnered with fact-checking organizations.

Regarding Media Literacy, LinkedIn does not prioritize any news sources in Feed; but in crisis situations, search banners may be used to point members to reputable sources of information.

Illegal Content and Activities

Risks related to “Illegal” Content and Activities include risk that “illegal” content and activities, such as intellectual property (IP) infringement or defamation; the promotion, concealment, acquisition, or sale of illegal, dangerous, or regulated goods (including

counterfeit goods, animals, drugs, psychoactive substances, criminal property, firearms, and other weapons) and services relevant to the same; the facilitation of illegal immigration; or activities causing the suffering of animals, are shared or perpetrated on the platform. Absent sufficient mitigations, risks related to “Illegal” Content and Activities may manifest on the platform in ways such as advertisers posting advertisements for illegal substances, members facilitating sale of illegal products via messages, members posting jobs that facilitate illegal immigration, or members posting content that constitutes IP infringement.

The **Inherent Probability** of systemic risks concerning “Illegal” Content and Activities stemming from the function, use, or misuse of LinkedIn’s services is assessed as “Not Likely.” Due to the professional nature of LinkedIn, it is “Not Likely” that members would publicly engage in illegal activities on the LinkedIn platform. The LinkedIn platform also does not allow for rapid and wide dissemination of content that has not been selected for quality, which limits systemic spread of illegal content. Considering the relative frequency of content or activities related to “Illegal” Content and Activities appearing in LinkedIn’s content moderation systems as well as public discourse related to LinkedIn, the probability is assessed as “Not Likely” that users would share or access “Illegal” Content or perpetrate Illegal Activities on LinkedIn’s platform absent sufficient mitigations.

The **Inherent Severity** of impact for content and activities that negatively impact “Illegal” Content and Activities on LinkedIn is rated as “Critical” primarily due to the gravity of risks within this category considering the potential for harm to security, environment, and wellbeing systems up to the regional and even the global level for environmental systems. Some risks within this category have potentially irremediable consequences.

With the rating for Inherent Probability as “Not Likely” and Inherent Severity as “Critical,” the **Inherent Risk** associated with negative impacts to “Illegal” Content and Activities across LinkedIn products and features is therefore assessed as “Moderate.”

In addition to the DTSP Best Practices LinkedIn has effectively implemented across all Risk Areas, which are further described under “[Mitigations](#),” LinkedIn has implemented robust mitigations specific to “Illegal” Content and Activities following best practices with defined, documented, and managed processes. As such, the **maturity of mitigations** LinkedIn has applied to the “Illegal” Content and Activities risk area is assessed as “Managed” which brings the **Residual Risk rating** down to “Low.”

Considering the Critical nature of this risk area, LinkedIn continued to invest in developing and enhancing strategies to further mitigate and manage risks related to “Illegal” Content and Activities, aiming for an even safer user experience. The key mitigations currently implemented specific to “Illegal” Content and Activities are described in the subsequent paragraphs of this section.

LinkedIn has two primary pathways to receive reports about content that may be illegal:

- The in-product reporting flow (available from the three-dot or “More” menu), which invites users to report content for a review of whether it complies with LinkedIn’s Professional Community Policies; and
- LinkedIn’s Illegal Content Reporting Form.

Reported content is queued for review and is removed if found to be in violation of LinkedIn’s policies or, if reported as illegal, applicable law. Reporting is available for all user-generated content on LinkedIn. User reports are instrumental not only in allowing LinkedIn to action policy-violating content where LinkedIn’s automated defenses did not detect it, but also in training these automated defenses to more accurately detect and timely enforce LinkedIn’s policies going forward.

LinkedIn’s User Agreement and Professional Community Policies specifically prohibit the violation of IP rights, including copyright, patents, trademarks, trade secrets or other proprietary rights. For example, users may not copy or distribute (except through available sharing functionality) the posts or other content of others without their permission. LinkedIn allows trademark owners to file claims per the process outlined in the Help Center. Similarly, LinkedIn allows copyright owners to file claims per the process outlined in the Help Center.

If LinkedIn becomes aware that potentially policy-violating content has been posted, LinkedIn reviews that content for policy compliance. LinkedIn works to proactively detect such content, and either remove it or route it for human review, using internally developed classifiers. These classifiers are developed with high levels of precision and recall, and are trained on a set of confirmed policy-violating content. In addition to this proactive review, members can flag potentially policy-violating content for human review by LinkedIn’s Trust Review Operations (TRO) team. Content review is conducted by a global team of language and policy enforcement experts who adhere to strict time limits in which the review must be completed. In addition, this team is subject to exacting training and quality assurance standards.

LinkedIn has a responsive approach to requests for removal of content that a user believes is defamatory or infringes IP rights. Using dedicated online forms, users can submit detailed complaints. LinkedIn processes and investigates each claim, and takes, documents, and communicates the appropriate action to the affected user(s). LinkedIn allows for defamation reports by the affected party per the process outlined in the Help Center.

If a member shares abusive, potentially illegal, or otherwise inappropriate content within the LinkedIn Feed, this content may not be eligible for broad distribution on the platform. Additionally, if other members are expressing disinterest in, dismissing or reporting a certain piece of content, that content may not receive broad distribution on the platform,

as this indicates members potentially find the content to be unhelpful, unwanted, or inappropriate for LinkedIn.

LinkedIn's Professional Community Policies apply to all members. Repeated offenses may result in feature restrictions. For certain egregious violations of LinkedIn's Professional Community Policies, LinkedIn may permanently restrict an account after a single violation. Users can leverage the searchable Help Center for additional details about LinkedIn policies and how they are enforced, including how to report abusive content. LinkedIn monitors edge cases for sensitive policy escalations to document decisions and recommend policy adjustments as needed.

Private and Family Life

Risks to Private and Family Life include malicious sharing or exploitation of sensitive or private data, doxing, or sharing of non-consensual private images. Absent sufficient mitigations, risks related to Private and Family life may manifest on the platform in ways such as members posting or threatening to post nonconsensual private images or information in Messages or members posting private information about another member or individual in Feed.

The **Inherent Probability** of systemic risks concerning Private and Family Life stemming from the function, use, or misuse of LinkedIn's services is assessed as "Not Likely." While LinkedIn limits the demographic or sensitive data collected by the platform and members tend to self-limit personal and family information shared on the platform, some members may share personal information on LinkedIn that may be exploited by bad actors and bad actors may share personal information about individuals from other sources to the platform. However, there is a lower likelihood of members sharing non-consensual intimate imagery or content on LinkedIn due to its professional nature. The LinkedIn platform also does not allow for rapid and wide dissemination of content that has not been selected for quality, which limits systemic spread of harmful content. Considering the relative frequency of content or activities related to Private and Family life appearing in LinkedIn's content moderation systems or in public discourse related to LinkedIn, the probability is assessed as "Not Likely" that users would share or access information or perpetrate activities negatively impacting Private and Family life on LinkedIn's platform absent sufficient mitigations.

The **Inherent Severity** of impact for content and activities that negatively impact Private and Family Life on LinkedIn is rated as "High" primarily due to the gravity of risks within this category, considering the potential for harm to economic, security, societal, and wellbeing systems at the individual level with some broader reverberations. Some risks in this category, such as posts that contain sensitive data or disclosure of nonconsensual private or intimate images, have potentially irremediable impacts. Some risks in this category, such as users' inadvertent sharing of sensitive personal

information (e.g. GDPR Article 9 information that includes racial or ethnic origin, political opinions, religious or philosophical beliefs etc.) may be less severe.

With the rating for Inherent Probability as “Not Likely” and Inherent Severity as “High,” the **Inherent Risk** associated with negative impacts to Private and Family Life across LinkedIn features is “Moderate.”

In addition to the DTSP Best Practices LinkedIn has effectively implemented across all Risk Areas, which are further described under “[Mitigations](#),” LinkedIn has implemented mitigations specific to Private and Family Life following industry best practices with defined, documented, and managed processes. As such, the **maturity of mitigation** efforts LinkedIn has applied to the Private and Family Life risk area is assessed as “Managed,” which brings the **Residual Risk** rating down to “Low.”

LinkedIn continues to invest in developing and enhancing strategies to further mitigate and manage risks related to Private and Family Life, aiming for an even safer user experience. The key mitigations currently implemented specific to Private and Family Life are described in the subsequent paragraphs of this section.

LinkedIn’s Professional Community Policies prohibit the posting of harassing content, including targeted personal attacks, intimidation, shaming, or disparagement, as well as revealing others’ personal or sensitive information (aka “doxing”) or inciting other to do the same. LinkedIn policies also prohibit sharing materials depicting nudity or sexual activity, creating fake profiles, and falsifying information about yourself or others.

If LinkedIn becomes aware that potentially policy-violating content has been posted, LinkedIn reviews that content for policy compliance. LinkedIn works to proactively detect such content, and either remove it or route it for human review, using internally developed classifiers. These classifiers are developed with high levels of precision and recall, and are trained on a set of confirmed policy-violating content. In addition to this proactive review, members can flag potentially policy-violating content for human review by LinkedIn’s Trust Review Operations (TRO) team. Content review is conducted by a global team of language and policy enforcement experts who adhere to strict time limits in which the review must be completed. In addition, this team is subject to exacting training and quality assurance standards. Additionally, if other members are expressing disinterest in, dismissing or reporting a certain piece of content, that content may not receive broad distribution on the platform, as this indicates members potentially find the content to be unhelpful, unwanted, or inappropriate for LinkedIn.

To help prevent fraudulent and malicious actors from accessing LinkedIn or from posting malicious content or scraping Profile data, LinkedIn has built automated detection systems to prevent bulk account creation attempts, identify the presence of a single bad actor creating or controlling multiple accounts and identify anomalous activity that may be indicative of a fake account.

To ensure that safety, security, and privacy remain at the core of the design and modification of LinkedIn's products, LinkedIn maintains a comprehensive pre-launch product review process. Among other things, this review considers how the product would collect and use data, to help LinkedIn measure privacy risk and ensure that appropriate data protection controls and safety measures are included.

Finally, LinkedIn helps protect members by empowering them to control the privacy of their own data, determining whether they can appear in search results on external search engines, how much of their Profile appears to LinkedIn members who are not first connections, and who can see their contact information. For European members who have affirmatively enabled certain enhanced safety features (through their settings), LinkedIn also displays safety reminders to such members if attempts to move private conversations away from the platform are detected in their incoming messages. The member can click through to view or report the message. This flow is designed to empower and educate members to make informed safety choices before they move certain conversations off platform.

Risk Areas of Minimal Residual Risk

Considering the probability, severity, and maturity of mitigations, LinkedIn assessed that two of the assessed risk areas fall into a Residual Risk category of “Minimal.” Public Security and Rights and Protection of Minors.

Public Security

Risk to Public Security includes risk that content or activities degrading public security occur on the platform including terrorist recruitment, funding, or training activities, terrorist imagery or content, or support for terrorists or glorification of terrorist acts as well as mis- or disinformation related to crisis events. Absent sufficient mitigations, risks related to Public Security may manifest on the platform in ways such as the members posted content glorifying terrorist acts in Feed, members posting a job that facilitates terrorist activities, or advertisers promoting content with misinformation on a public crisis event.

The **Inherent Probability** of systemic risks concerning Public Security stemming from the function, use, or misuse of LinkedIn’s services is assessed as “Remote.” While data analysis still shows some demand for posting and consuming content related to world crisis events, such as the Israel-Hamas war, it is lower than last year and there is a lower likelihood of members sharing such content on LinkedIn due to its professional nature. The LinkedIn platform also does not allow for rapid and wide dissemination of content that has not been selected for quality, which limits systemic spread of harmful content. Based on the relative frequency of content or activities related to Public Security appearing in LinkedIn’s content moderation systems as well as public discourse related to LinkedIn, the probability is assessed as “Remote” that users would share or access information or perpetrate activities negatively impacting Public Security on the LinkedIn platform absent sufficient mitigations.

The **Inherent Severity** for content and activities that negatively impact Public Security on LinkedIn is rated as “Critical,” primarily due to the gravity of risks within this category, considering the potential for significant harm to security, wellbeing, societal, political, and economic systems at the individual, country, and regional levels. The impact of many risks within this category are irremediable.

With the rating for Inherent Probability as “Remote” and Inherent Severity as “Critical,” the **Inherent Risk** associated with negative impacts to Public Security across LinkedIn products and features is therefore assessed as “Low.”

In addition to the DTSP Best Practices LinkedIn has effectively implemented across all Risk Areas, which are further described under [“Mitigations,”](#) LinkedIn has implemented robust mitigations specific to Public Security following industry best practices with defined and documented processes. As such, the **maturity of mitigation** efforts LinkedIn has applied to the Public Security risk area is assessed as “Managed,” which brings the **Residual Risk rating** down to “Minimal.”

Considering the Critical nature of risks in this category, LinkedIn continues to invest in the development and enhancement of strategies to further mitigate and manage risks related to Public Security, aiming for an even safer user experience. The key mitigations currently implemented specific to Public Security are described in the subsequent paragraphs of this section.

Upon registration, LinkedIn members agree to abide by the User Agreement and the Professional Community Policies, which expressly forbid members from posting information that promotes dangerous organizations or individuals, including terrorists. Specifically, the Professional Community Policies provide:

Do not share content promoting dangerous organizations or individuals. We do not allow any terrorist organizations or violent extremist groups on our platform. And we do not allow any individuals who affiliate with such organizations or groups to have a LinkedIn profile. Content that depicts terrorist activity, that is intended to recruit for terrorist organizations, or that threatens, promotes, or supports terrorism in any manner is not tolerated.

The policy also links out to a [Help Center article](#) containing examples of dangerous organizations and individuals content types prohibited by the policy. Additionally, LinkedIn restricts Profiles and Pages associated with dangerous organizations and individuals regardless of whether they have posted violative content.

LinkedIn's Professional Community Policies apply to all members. Repeated offenses may result in feature restrictions. For certain egregious violations of LinkedIn's Professional Community Policies (e.g., terrorism, extremely violent content), LinkedIn may permanently restrict an account after a single violation.

Content review is conducted by LinkedIn's Trust Review Operations (TRO) team. TRO is a global team of language and policy enforcement experts that adhere to strict time limits in which the review must be completed (SLAs). The time limits are dependent on the severity of member harm, and certain types of content, e.g., potential terrorist content, is human reviewed and removed as soon as possible. In addition, the TRO team is subject to exacting training and quality assurance standards, and LinkedIn has separate teams that effectuate these standards. LinkedIn's team of internal content policy professionals also provide specific civic and election-related guidance to TRO as appropriate given the circumstances.

LinkedIn also works closely with Global Internet Forum to Counter Terrorism (GIFCT) and other civil society groups to receive inputs to its content policies and additional information on how its platform can be abused to harm users. As a member of GIFCT, LinkedIn has access to GIFCT's Incident Response processes, including hashes related to an event activated as Content Incidents or Content Incident Protocols.

LinkedIn uses detection systems that identify the presence of a single bad actor creating or controlling multiple accounts, and it identifies anomalous activity that may be indicative of a fake account. LinkedIn's automatic detection filters for policy violating content prior to being posted on the platform. LinkedIn also relies on member reports and a team of investigators to uncover accounts and groups that may have evaded all other levels of defense. LinkedIn also publishes information regarding its responses to government removal requests and voluntarily discloses certain information to relevant law enforcement authorities, as permitted by law, in the event LinkedIn becomes aware of certain imminent threats.

Rights and Protection of Minors

Risks related to Rights and Protection of Minors include child sexual exploitation, grooming, or sharing Child Sexual Abuse Material (CSAM); collection of minor's data without parental consent; exposure of children to harmful content or activities; cyberbullying and harassment of children; or promotion of self-harm to minors. Absent sufficient mitigations, risks related to Rights and Protection of Minors may manifest on the platform in ways such as members attempting to contact minors through InMail for grooming purposes, members soliciting or sharing child exploitation material through Messages, advertisers posting ads that promote dangerous or harmful products or services to minors, or members posting harmful or harassing content directed at minors within Feed.

The **Inherent Probability** of systemic risks concerning Rights and Protection of Minors stemming from the function, use, or misuse of LinkedIn's services is assessed as "Remote." LinkedIn, as a specific-purpose professional network is not designed for or used heavily by users under the age of 18. There is also a lower likelihood of members sharing material related to child exploitation on LinkedIn due to its professional nature. LinkedIn has historically and consistently detected only limited instances of child sexual abuse material on the platform. Considering the relative frequency of content or activities related to the Rights and Protection of Minors appearing in LinkedIn's content moderation systems as well as public discourse related to LinkedIn, the probability is assessed as "Remote" that users could share or access information or perpetrate activities negatively impacting the Right and Protection of Minors on LinkedIn's platform absent sufficient mitigations.

The **Inherent Severity** for content and activities that negatively impact the Rights and Protection of Minors on LinkedIn is rated as "Critical" due to the gravity of risk in this category, the vulnerable nature of the population impacted, the irremediability of impact of certain risks, namely CSAM, and the potential for societal impact broader than the individual scale for risks impacting youth. Not all risks within this category are assessed at the highest level of severity, namely collection, processing, and use of children's data without parental consent.

With the rating for Inherent Probability as “Remote” and Inherent Severity as “Critical,” the **Inherent Risk** associated with negative impacts to the Rights and Protection of Minors across LinkedIn products and features is “Low.”

In addition to the DTSP Best Practices LinkedIn has effectively implemented across all Risk Areas, which are further described under “[Mitigations](#),” LinkedIn has implemented robust mitigations specific to the Rights and Protection of Minors following industry best practices with defined, documented, and managed processes. As such, the **maturity of mitigation** efforts LinkedIn has applied to the Rights and Protection of Minors risk area is assessed as “Managed,” which brings the **Residual Risk rating** for down to “Minimal.”

Considering the Critical nature of this risk area, LinkedIn continues to invest in developing and enhancing strategies to further mitigate and manage risks related to the Rights and Protection of Minors, aiming for an even safer user experience. The key mitigations currently implemented specific to the Rights and Protection of Minors are described in the subsequent paragraphs of this section.

LinkedIn is a professionally oriented platform. It exists to allow the world’s professionals to connect, learn, hire, be hired, and engage in other *professionally-oriented* activities. LinkedIn’s Professional Community Policies, which are incorporated by reference into its User Agreement, twice specify that the nature of communication on the platform must remain professional:

Be professional

We allow broad conversations about the world of work, but just keep it professional.

...

Be professional

LinkedIn's mission is to connect the world's professionals to make them more productive and successful. We allow broad conversations about the world of work, but require professional expression.

Given its professional focus, content, and context (content about the world of work, job leads and applications, professional development, etc.), LinkedIn is generally neither attractive to minors nor a platform that a significant number of minors use. Consistent with this professional environment, LinkedIn is also not designed in a manner that would be attractive to minors. For example, the color scheme used throughout the LinkedIn service is generally toned-down and professional, not bright and bold.

LinkedIn’s terms of service require that members must be at least 16 years old. Specifically, the LinkedIn User Agreement states:

To use the Services, you agree that: (1) you must be the "Minimum Age" (described below) or older...Creating an account with false information is a

violation of our terms, including accounts registered on behalf of others or persons under the age of 16.

“Minimum Age” means 16 years old. However, if law requires that you must be older in order for LinkedIn to lawfully provide the Services to you without parental consent (including using your personal data) then the Minimum Age is such older age.

Accordingly, only a subset of minors (i.e., 16- and 17-year-olds) are eligible to use LinkedIn. Would-be users who are identified as being under 16 at time of account creation are blocked from creating an account. If LinkedIn otherwise becomes aware that a user is under LinkedIn’s minimum age after account sign-up, their account is restricted. LinkedIn may retain the information necessary to prevent the minor from opening an account using the same email address, and other data associated with the account is purged unless retention is compelled under applicable law.

LinkedIn does facilitate targeted advertising that is in part based on profiling. It also does infer age based on objective criteria input by members on their profile, including graduation date. LinkedIn does not, however, infer an age under eighteen, and does not enable ad targeting to age groups below eighteen.

While LinkedIn has no age inferences for members younger than eighteen and therefore would not “be aware with reasonable certainty that the recipient of the service [receiving an ad] is a minor” (see Art. 28(2)), we are still taking additional steps out of an abundance of caution to ensure that profiling ads are not received by the limited number of members who may be under 18. To do this, members in the European Union who are listed as attending a high school or secondary education institution on their profile are excluded from all targeted advertising. They may still see ads, but they will not be personalized based on any profiling or activity information.

Although most minors are prohibited from becoming LinkedIn members and those on the platform are not many, LinkedIn’s policies generally prohibit the types of content that would be harmful to minor users, including for example content containing pornography, promotion of self-harm or suicide, bullying, unsafe challenges, terrorist content, and child exploitation material. As a result, and as described throughout this Report, LinkedIn uses a range of proactive reactive measures to detect and remove this type of content from LinkedIn so users of all ages do not encounter it.

Likewise, LinkedIn offers its members of all ages—including minors—a range of user empowerment tools to help them control what content they see, how their information is viewed by others, and how others interact with them. For example, members can:

- Control how their Profile is shown on LinkedIn as well as on search engines.

- Control who can send them an invitation to connect on LinkedIn (e.g., members can choose to receive connection invitations only from people who already know their email address);
- Disallow non-connections from sending them messages (to avoid unwanted contact);
- Block particular members by using LinkedIn's "block" feature;
- Restrict the visibility of their email address, list of connections, and last name; and
- Disallow others from being able to find their LinkedIn Profile using the member's email address or phone number.

LinkedIn has developed advanced detection capabilities related to harmful content that underage users might be exposed to on LinkedIn, including illegal or dangerous goods or services, dangerous organizations or individuals, and hateful, harmful, shocking, violent, or harassing content. Specifically, for CSAM, LinkedIn scans all public-facing and privately shared videos and images for CSAM (including content that is nearly identical to confirmed CSAM content), conducting hash matching against known CSAM imagery. Detected CSAM is removed or prevented from being shared on the platform, reported to the National Center for Missing and Exploited Children (NCMEC), preserved to enable law enforcement investigations, and the associated account is placed in restriction. Once CSAM has been identified, LinkedIn also applies algorithmic models to identify and remove near-duplicates from the platform. Although the numbers of confirmed CSAM cases on LinkedIn are quite small, as explained in LinkedIn's Transparency Report, LinkedIn remains diligent and focused on this anti-abuse work.

LinkedIn currently is reviewing the July 2025 Guidelines published by the European Commission in connection with Article 28 of the DSA to determine which mitigations recommended therein are proportional and appropriate to ensure the safety, security, and privacy of minors on LinkedIn. Based on the findings of that review, LinkedIn may refine or supplement its existing mitigation measures for this risk area during the next risk assessment period.

6.2 Scoring Rationale

Probability Across Risk Areas

The Inherent Probability or likelihood of occurrence for most systemic risks is assessed at relatively low levels. The professional nature of the platform discourages typical members from posting potentially harmful content that would be visible to their employers and potential or current business partners. This determination is supported by the fact that the volume of violative content on LinkedIn is consistently lower than industry norms, when compared with industry Transparency Reports.

Considering (1) content flagged and removed from the platform, by automated classifiers and human reporting, (2) the volume of public discourse related to LinkedIn and each risk area, and (3) trends in public concern and emerging threats across the industry, LinkedIn assessed Consumer Protection and Fraud as having the highest relative probability of occurrence on the platform, followed by Human Dignity, Mental and Physical Wellbeing, Discrimination and Hate, Civic Discourse and Electoral Processes and Freedom of Expression and Information. Public Security and the Rights and Protection of Minors were assessed at the lowest levels of probability.

Impact of Influencing Factors

In the assessment of each risk area and its associated probability, severity, and mitigation effectiveness, LinkedIn considered the potential impact of all Article 34 influencing factors, including the design of recommender systems and algorithmic systems; content moderation systems; applicable terms and conditions and their enforcement; systems for selecting and presenting advertisements; LinkedIn data related practices; intentional manipulation of the service; amplification or rapid and wide dissemination; and specific regional or linguistic aspects.

Recommender Systems and Rapid Dissemination: LinkedIn strives to broadly distribute only high-quality content – content that meets specific quality and professionalism criteria – in Feed. In practice, this means content that shares knowledge or advice on professional topics that LinkedIn thinks its members will likely be interested in will get broader distribution, whereas divisive, unprofessional, or harmful content should not be amplified or broadly disseminated. LinkedIn also provides suggested content that is valuable and relevant to those outside of the original poster’s network. Finally, the GenAI features that LinkedIn provides can only be used in limited, context-specific scenarios, thereby limiting the widespread use of AI systems capable of generating text, images, videos, or other synthetic content and the dissemination of such content on its platform. This aligns with LinkedIn’s view that the goal of LinkedIn is not virality or reach at all costs; instead, it’s about reaching the right professionals with the right information they care about that can help in their careers.

Algorithmic Systems: LinkedIn has a [robust process](#) in place to examine the fairness and potential bias of algorithmic systems deployed in any new feature or product. Algorithmic review is included as part of LinkedIn's Trustworthy Design Review (TDR) process and regular fairness checks of Feed and content distribution. LinkedIn is committed to upholding trust, promoting fairness and inclusion, providing transparency, and embracing accountability with the use of AI on its platform.

Content Moderation: LinkedIn's approach to content moderation is risk-based and carefully weighs safety against freedom of expression, erring toward freedom of expression whenever possible. LinkedIn employs advanced detection techniques and sophisticated algorithms to detect [policy-violating content](#) and activity on the platform before it causes harm. LinkedIn uses a layered, multidimensional approach to moderate content within LinkedIn's Trust ecosystem, including automatic and proactive prevention as well as human-led detection. In some cases, LinkedIn may permanently restrict a member's access to certain LinkedIn features after multiple policy violations or suspend an account in response to certain single egregious content policy violations (e.g., in the case of child exploitation material).

Final decisions on content removal or account suspension made by LinkedIn's human reviewers feed back into the system to further train the platform's models. LinkedIn measures its preventive defense services regularly to improve accuracy in the filtering process.

Advertisements: It is important to note that a majority of LinkedIn's advertisements are business-to-business rather than business-to-consumer, which has a significant impact on LinkedIn's risk profile related to advertising systems. LinkedIn maintains and enforces [policies](#) related to advertising on the platform. These policies include, among other things, prohibitions on political advertisements; policies that limit healthcare-related advertisements to come from certified retailers and that limit distribution to jurisdictions where such advertisements are legal; and policies that prohibit discrimination based on protected characteristics. LinkedIn has also updated its advertising tools to remove the ability to target members in Europe using LinkedIn Groups, to prevent any misconception that ads could be targeted at members using special categories of data. LinkedIn also allows members in Europe to choose whether to connect their data across LinkedIn features and solutions.

Intentional Manipulation of the Service: As a real-identity platform, LinkedIn addresses intentional manipulation at the source by investing significantly to prevent the creation of [fake accounts](#). LinkedIn's teams have created and deployed sophisticated models to block bulk account creation, identify clusters of bad accounts, detect anomalous account behavior, identify AI-generated Profile photos, and remove phishing links and malware. In the last Transparency Reporting period, LinkedIn's automated defenses blocked 97.1% of the fake accounts that were caught during the July -

December 2024 period, with the remaining 2.9% stopped by manual investigations and restrictions. 99.7% of the fake accounts were stopped proactively, before a member report.

Linguistic and Regional Considerations: LinkedIn currently supports the following languages on its platform: English, Arabic, Bangala, Chinese (Simplified and Traditional), Czech, Danish, Dutch, Finnish, French, German, Greek, Hebrew, Hindi, Hungarian, Indonesian, Italian, Japanese, Korean, Malay, Marathi, Norwegian, Persian, Polish, Portuguese, Punjabi, Romanian, Russian, Spanish, Swedish, Tagalog, Telugu, Thai, Turkish, Ukrainian, and Vietnamese. LinkedIn has enabled automated and user-led detection of policy-violating content in all supported languages. Regarding European elections, LinkedIn engages with stakeholders as needed to address rising concerns relative to election cycles and ensure appropriate linguistic and cultural coverage.

Year 3 Changes in Risk Scores

The Residual Risk scores for only two risk areas changed from the 2024 to the 2025 assessment, as shown in Figure 34. Some Inherent Probability and Maturity of Mitigations scores shifted this year as well. However, those changes did not impact LinkedIn's Residual Risk Scores. Key changes are described here.

Figure 34: Changes In Residual Risk Scores

Risk Area	Residual Risk Rating
Consumer Protection and Fraud	Moderate
Human Dignity	Moderate
Mental and Physical Wellbeing	Low
Discrimination and Hate	Low
Freedom of Expression and Information	Low
Protection of Personal Data	Low
Civic Discourse and Electoral Processes	Low
Public Health	Low
Illegal Content and Activities	Low
Private and Family Life	Low +
Public Security	Minimal -
Rights and Protection of Minors	Minimal

Key: + indicates an increase of a rating from last year's assessment by one rating level
- indicates a decrease of a rating by one rating level

- **The Inherent Probability of risks to Private and Family Life increased from “Remote” to “Not Likely” and therefore Residual Risk score increased from “Minimal” to Low.”** Due to increased volumes of discussion on social media about online privacy concerns (particularly those relating to GenAI enhancements and the creation and online sharing of non-consensual intimate imagery), Inherent Probability for this risk area rose from “Remote” last year to “Not Likely” this year. This increased the Inherent Risk Rating from “Low” to “Moderate,” driving a higher Residual Risk this year at “Low.”
- **The Inherent Probability of risks to Public Security decreased from “Not Likely” to “Remote,” and Maturity of Mitigations increased from “Defined” to “Managed”; therefore the Residual Risk score decreased from “Low” to “Minimal.”** Due to lower volumes of public discourse related to terrorism compared to the prior assessment period (as evidenced by internal and external data), the Inherent Probability for this risk area decreased from “Not Likely” last year to “Remote” this year. In addition, as part of a deeper review of mitigations related to terrorist content, LinkedIn identified additional measures implemented to combat risks related to Public Security, resulting in an increase in the rating of the Maturity of Mitigations from “Defined” to “Managed.” Due to the decrease in Inherent Risk and the increase in Maturity of Mitigations, the Residual Risk score reduced from “Low” last year to “Minimal” this year.
- **The Inherent Probability of risks related to Civic Discourse and Electoral Processes decreased from “Likely” to “Not Likely” while Residual Risk remained “Low.”** In part due to the decrease in the total number of elections around the world from last year to this year, the volume of discussion on social media concerning elections was lower. Despite this change, the Residual Risk Rating remains at “Low.”

7. Mitigations

DTSP participating companies commit to five fundamental areas of best practices to support the operation of trusted and safe products and services: (1) product development, (2) governance, (3) enforcement, (4) improvement, and (5) transparency. Underpinning these are the 35 specific best practices of the DTSP Best Practices Framework, which provide concrete examples of different activities and processes that companies may implement to mitigate risks associated with harmful content and conduct. By committing to these best practices, LinkedIn has aligned with industry-accepted Trust and Safety standards, enabling it to better identify and manage potential risks. By mapping these mitigations to each of the DTSP Best Practice Buckets, LinkedIn creates greater transparency and provides reliable and effective solutions to challenges associated with harmful content and conduct.

7.1 Best Practice Highlights

Product Development

Trustworthy Design Review (TDR): The Trustworthy Design Review (TDR) process is one of LinkedIn's primary tools for risk identification and mitigation during the product development lifecycle. TDR helps LinkedIn's features meet LinkedIn's high standards for protecting members' privacy and safety. TDR also serves to loop in and centralize communications between relevant stakeholders from across LinkedIn's Product, Legal, Trust Review Ops and InfoSec teams. At LinkedIn, new products and initiatives undergo a pre-launch review and assessment by cross-disciplinary specialists, with a focus on legal compliance, data protection, security, and safety. The teams first exchange certain key information about the new product or initiative to determine if a privacy impact assessment, security review, and/or safety assessment should be conducted as part of the TDR process. This means that Trust and Safety consultation is built into the review process for new products and initiatives. During this evaluation, LinkedIn reviews and documents any potential risks and abuse vectors associated with new products or features including, but not limited to, considerations of member trust, necessary tooling, policy, support requirements, member education, transparency, public relations risks, training needs, and trust design recommendations, along with their respective mitigations. The process provides feedback to enhance the product or initiative and improve relevant Trust and Safety measures and helps to ensure compliance with data protection commitments, security standards, and applicable laws, as well as alignment with LinkedIn's Guiding Principles for Trust. As part of TDR, LinkedIn's legal team also reviews any proposed new GenAI feature on the platform, including to ensure members are aware that a feature uses GenAI and – if GenAI content is suggested to the member – the member should review and revise such content, as they're responsible for what they ultimately publish or send.

Abuse Pattern Analysis: LinkedIn has a dedicated abuse prevention team that prioritizes, triages, and remediates numerous abuse vectors (both acute and systemic) in order to keep LinkedIn members safe. LinkedIn's Trust Product team orchestrates regular prevalence exercises where it manually reviews samples of accounts and content on platform to identify new patterns of abuse on the site, which inform both design of future defenses and priorities in its Trust and Safety roadmap. The Trust Product team has also implemented several proactive defenses to detect fraudulent behavior and policy-violating content. LinkedIn monitors the performance of these defenses regularly to ensure they are working optimally and address any gaps as they are identified.

Trust and Safety Consultation: In addition to the TDR process, LinkedIn's Trust Review Ops (TRO), Trust Engineering, Trust Product, and Trust Data Science teams work in concert through the entire product development, strategic planning, and budgeting lifecycles to ensure alignment of LinkedIn's priorities and initiatives with Trust Principles.

Accountability: LinkedIn has built accountability into the TDR process by managing cross-functional reviews in a centralized tracking and project management system to help ensure accountability and thoroughness of the review process.

Feature Evaluation: As described above, specific features are assessed holistically as part of the TDR process before launch. For algorithmic features, LinkedIn has also developed tools to enable the measurement of fairness in large-scale machine learning and artificial intelligence workflows. It can be deployed in training and scoring workflows to measure biases in training data, evaluate different fairness notions for ML models, and detect statistically significant differences in model performance across different subgroups.

Specifically, when building out GenAI features, Trust, InfoSec and Legal teams provide detailed list of requirements (e.g., information on Responsible AI Transparency and Labeling requirements) for product teams to abide by during the TDR process. Together, these teams work to comply with the law, ensure member trust, comply with LinkedIn's Responsible AI Principles, and ensure LinkedIn's applications, features, and data are secure.

LinkedIn is mindful that GenAI models can reflect the stereotypes that exist in society and actively works to avoid harmful stereotypes in GenAI tool outputs by deploying multi-layered tactics throughout the product development lifecycle. LinkedIn conducts fairness checks on material changes to the algorithms for LinkedIn Feed and LinkedIn Job recommendations, conducts A/B testing, and holds out groups of members to assess impacts of algorithmic feature launch. For instance, to mitigate the risk of abuse on job descriptions (for instance through injection of harmful text into the prompt) in LinkedIn Jobs, the team has deployed a drop-down menu that includes standardized job

titles. This measure enables the generation of job descriptions that are responsive to the customer's input while also reducing the potential for abuse and bias. Similarly, with personalized writing suggestions for your Profile (while it's still based on information already on the member's Profile) LinkedIn constructed the prompts in a manner that reduces the risk of problematic output (resulting from "jailbreaking" the AI tool through nefarious prompts that aim to end run restrictions in AI that are designed to maximize the chance that outputs will be reasonable and appropriate) and the member feedback menu includes the option to flag "biased content."

Finally, LinkedIn works hard to ensure that its products meet the W3C (World Wide Web Consortium) WCAG (Web Content Accessibility Guidelines) V2.1 Level AA. As part of its commitment to accessibility LinkedIn continuously audits its products—internally and through a third party—using assistive technology such as screen reading software (e.g., JAWS, NVDA, VoiceOver and TalkBack). LinkedIn has added reflow and orientation support to the LinkedIn app, which means that not only its policies, but also the reporting flow can be accessed in a vertical or horizontal orientation, which is helpful to members with certain disabilities. Likewise, LinkedIn also allows for dark mode and adjust of font size, both of which are reflected on its reporting mechanism. If users need help, LinkedIn's Disability Answer Desk assists members who use assistive technology with LinkedIn-related issues.

Risk Assessment: LinkedIn's Digital Safety function, partners with the relevant product and engineering organizations to conduct continuous assessments of risk during the full product development lifecycle. This ensures a "safety by design" approach throughout the process. Furthermore, LinkedIn's Compliance Team performs LinkedIn Compliance deep dives and assists with the Company's Internal Audit workstreams. LinkedIn also has a robust Trust Product planning process (including the Trustworthy Design Review or "TDR" process described above) to address emerging threats and risks.

Pre-Launch Feedback: Pre-launch feedback from cross-functional stakeholders is incorporated as part of the TDR review process. This process enables LinkedIn to evaluate new products and features for risk and abuse vectors, including but not limited to member trust, tooling needs, policy, support requirements, communication, transparency, PR risks, potential abuse vectors, training needs and trust design recommendations. For GenAI features, LinkedIn has measures in place to help ensure that any model used, and feature launched is compliant with LinkedIn's Responsible AI Principles. Model alignment focuses on fundamental risks (e.g., sexual content) and not on product specific risk (e.g., political content). User feedback metrics are also used to serve as the guardrails for and to improve all GenAI features and models.

Post-Launch Evaluation: In addition to the TDR process, the product team(s) responsible for the relevant product or initiative lead ongoing engagement with the other trust stakeholders from the TDR process to learn from real risk patterns and implement any additional safety features post-launch. This process may include fairness checks and other quality assurance activities further described under Effectiveness Testing. If a high-risk situation is identified post-launch, LinkedIn teams take mitigating measures before the product or feature is fully ramped.

User Feedback: LinkedIn has an active external and internal communications team that interviews users across a variety of topics and cascades member feedback to various internal teams. LinkedIn has a “feedback” button on the site for users to provide feedback. Users can also directly provide feedback about Trust issues using several in-product features like “Reporting,” “Blocking,” “I don’t know this person,” or by reaching out to LinkedIn’s customer support teams. This feedback is incorporated via various automated and manual means into LinkedIn’s proactive defenses to help improve accuracy and coverage. LinkedIn monitors metrics on these reports on a regular basis to identify and address problematic trends. These learnings are then used to inform its Trust roadmap. In addition, LinkedIn’s User Experience and Research (UXR) and Market Research teams partner with product and design stakeholders to gather feedback directly from LinkedIn members, customers, and internal users of Trust tools throughout the innovation process; foundational research feeds ideation based on member Trust experiences and evaluative usability research on proposed Trust features.

User Controls: LinkedIn provides platform and feature specific controls that allow members to tailor their experience on the platform via user controls and settings.

- **Cross-Feature Controls:**
 - Members in the EEA and Switzerland can choose whether to connect their core LinkedIn professional networking experience with our Jobs, Marketing Solutions and/or Learning services. When connected LinkedIn shares data from the core LinkedIn experience with these other services to provide more tailored experiences, and to help improve these services.
 - Members can decide if they want to view content labeled as sensitive by LinkedIn or not. LinkedIn will also warn members when photos or videos that may contain sensitive or graphic content are detected. The content will initially be blurred with a warning, but members can decide if they would like to view it.
- **Feed:** In the Feed, a “Sort by” feature is available in the user interface, allowing members to change their preferred sorting option for their Feed, thus providing an alternative to the recommender system. The “Most recent” setting arranges updates by publication date, while the “Most relevant” setting ranks them by relevance leveraging the Feed Recommender System. LinkedIn’s “I don’t want to

see this” option in the overall Feed surface allows members to prevent specific posts from appearing in their Feed again. Members also have granular control over notifications through their settings. Members can also express their disinterest in certain topics, e.g., hiding content in their Feed.

- **Ads:** All LinkedIn members have the option to select within their settings whether they would like their Profile information to be used to personalize the ads visible to them on LinkedIn.
- **Games:** Configurable settings exist to turn on/off the sharing of Games-related content. Games scores are private unless a member chooses to share them.
- **Messages:** LinkedIn offers settings that allow members to only receive communications from people they know.
- **Profile:** Members can manage the visibility of their information through settings and also who their data is shared with (i.e., post visibility). These settings are honored throughout the LinkedIn platform, including on Feed, Search, and other personalized modules like People You May Know and Notifications. For example, members can:
 - Control how their Profile is shown on LinkedIn, such as on search engines.
 - Control who can send them an invitation to connect on LinkedIn (e.g., members can choose to receive connection invitations only from people who already know their email address);
 - Disallow non-connections from sending them messages (to avoid unwanted contact);
 - Block particular members by using the “block” feature (see more here);
 - Restrict the visibility of their email address, list of connections, and last name; and
 - Disallow others from being able to find their LinkedIn Profile using the member’s email address or phone number.
 - Turn on/off the sharing of Games-related content. Games scores are private unless a member chooses to share them.
 - LinkedIn allows users to turn off comments on their posts, as explained in the Help Center article titled “Turn off and Limit Comments on Posts.”
 - In addition, LinkedIn enables members to control who can find them (via invitations to connect, job opportunities, etc.), allowing them to make only those connections that are meaningful.

Product Governance

Policies & Standards: LinkedIn's Content Policy team, which collaborates closely with other parts of Legal as well as the Trust and Safety, Product, and Engineering teams, maintains detailed internal policies and procedures related to enforcement of the [Professional Community Policies](#). LinkedIn's terms of service are updated periodically, as needed. Material updates to the User Agreement and Privacy Policy are communicated directly to members – by means of emails and in-product notifications in advance of the update, together with a preview of the update – and LinkedIn provides persistent links that allow users to access the terms.

Jobs

Job posts on LinkedIn must also comply with LinkedIn's [Job Policies](#). Additionally, with respect to job risks, LinkedIn's Jobs Policies also clearly specify that fraudulent job posts are not allowed on LinkedIn. Relevant provisions of the Jobs Policies, which LinkedIn specifically enforces, include:

- Job posts must be legitimate. Only create postings for bona fide job opportunities with a reasonable and legitimate intent to hire for the specific position listed. Job posts created primarily for user acquisition (e.g., directing job seekers to sign up for an external platform or service), self-promotion, or purpose other than hiring for a genuine job opportunity are not permitted. Each job post must be for one job opportunity; you cannot solicit general applications for roles at a company.
- Fraud and Deception. Job posts must not be fraudulent or deceptive. Do not intentionally misrepresent the job, hiring company, or poster. Do not create job posts that require applicants to pay in order to apply or for employment. Failure to clearly disclose in any posting that a position is for an independent contractor or is part-time, piecework, commission-based, or has otherwise non-traditional working conditions or compensation may constitute intentional misrepresentation.
- MLM and Franchises. Do not create postings for "business opportunities" that require payments or recruitment of others or that resemble franchises, multi-level marketing, club memberships, or distributorships.
- Malicious Software. Do not create postings that contain malware, spyware or any other malicious code or otherwise interfere with the operation of the Jobs Services or any device or system or breach or circumvent any security measure of LinkedIn or a third party.
- Phishing. Do not use a job posting to promote a website that tricks a user into providing personal or other information.

LinkedIn has repeated offender systems to limit the ability of bad actors to repeatedly post jobs that violate LinkedIn's policies (including fraudulent jobs). LinkedIn has implemented systems at both the job poster and the LinkedIn Page level. That is, a job poster that posts multiple jobs (or in some cases a single job, depending on the policy

violation) that violates LinkedIn's policies will have its job posting ability restricted. Similarly, a LinkedIn Page with job posts that repeatedly violate LinkedIn's policies will have its job posting ability restricted. This helps limit bad actors from attempting to conduct fraudulent schemes at scale.

Minors

Only a subset of minors (i.e., 16- and 17-year-olds) are eligible to use LinkedIn. Would-be users who are identified as being under 16 at time of account creation are blocked from creating an account. If LinkedIn otherwise becomes aware that a user is under LinkedIn's minimum age after account sign-up, their account is restricted. LinkedIn may retain the information necessary to prevent the minor from opening an account using the same email address, and other data associated with the account is purged unless retention is compelled under applicable law.

Advertising

Furthermore, for online advertisers, LinkedIn sets forth detailed policies and standards in the LinkedIn Ads Agreement regarding launching ad campaigns and all online advertisers are required to abide by the requirements listed to launch successful ad campaigns. LinkedIn also has global ad policies in place that do not allow targeting sensitive groups on the platform. This decision helps to protect LinkedIn members and customers.

With respect to advertising risks, LinkedIn does not accept political ads and has not done so since 2018. LinkedIn prohibits political advertisements on its platform, including ads advocating for or against a candidate, party, or ballot proposition or ads designed for political fundraising or otherwise intended to influence an election outcome. LinkedIn maintains a public Ad Library that allows members, researchers, the media, and the public to monitor ads that have run on LinkedIn, as well as LinkedIn's enforcement of its Advertising Policies. With respect to the risk of hyper-targeted political ads, as an initial matter, as mentioned above, LinkedIn does not allow political ads and has implemented systems and processes to reject them. In addition, LinkedIn mitigates the risk of hypertargeting advertisements (across all ads, including attempted political ads) by implementing minimum audience size requirements before a campaign can run, preventing an advertisement from being hyper targeted to a specific individual.

Accessibility

Finally, LinkedIn's User Agreement is written in plain, comprehensible language, and meets the W3C (World Wide Web Consortium) WCAG (Web Content Accessibility Guidelines) V2.1 Level AA. As part of its commitment to accessibility LinkedIn continuously audits its products—internally and through a third party—using assistive technology such as screen reading software (e.g., JAWS, NVDA, VoiceOver and TalkBack). Reader comprehension is impaired when a single document is too lengthy or wordy, so LinkedIn improves accessibility by pulling some information out into separate

Help Center articles, which are incorporated into the User Agreement and/or Professional Community Policies. Links to Help Center articles providing more information on topics are clearly signposted using blue underlined text. LinkedIn holds itself accountable through transparency and periodically publishes Accessibility Conformance Reports (ACR's) at linkedin.com/accessibility/reports that report on the accessibility of its products.

Changes

LinkedIn revises its policies (or develops new ones) as appropriate when changes in the regulatory environment or the needs of the business, including new trends related to illegal harms on the service, require us to clarify or rethink our approach to content moderation. Product ramps introduce new content sources through which policy violations manifest in different ways, and require additional guidance to consistently enforce.

User Focused Product Management: In its Feed, LinkedIn strives to broadly distribute only high-quality content that meets specific quality and professionalism criteria. In practice, this means content that shares knowledge or advice on professional topics that LinkedIn thinks its members will likely be interested in will get broader distribution, whereas divisive, unprofessional, or harmful content will not be amplified or broadly disseminated. This aligns with LinkedIn's view that the goal of LinkedIn is not virality or reach at all costs; instead, it is about reaching the right professionals with the right information they care about that can help in their careers. LinkedIn leverages user feedback to limit the visibility of content that is less likely to lead to an enhanced professional experience.

Community Guidelines/Rules: LinkedIn's public-facing policies are set forth in the User Agreement, Privacy Policy, Professional Community Policies, Jobs Terms & Conditions, and Ads Policies, each of which (1) can easily be accessed on the LinkedIn site and within the product, and (2) is publicly available in a number of languages. LinkedIn's User Agreement and Privacy Policy also are presented to members upon registration. Users can leverage LinkedIn's searchable Help Center for additional details about its policies and how LinkedIn enforces them, including how to report abusive content. LinkedIn's policies include provisions outlining what is and is not allowed on the platform. The Professional Community Policies apply to all members and content.

On joining LinkedIn, members agree to abide by LinkedIn's User Agreement and Professional Community Policies. The Professional Community Policies are incorporated by reference and form part of LinkedIn's User Agreement. The Professional Community Policies also, in turn, incorporate a number of links to Help Center articles explaining how LinkedIn minimizes the length of time for which policy-violating content is present on the service, including information about the use of

proactive technologies and information regarding the policies and processes that govern the handling and resolution of relevant complaints.

LinkedIn Professional Community Policies prohibit a range of abusive content, including harassing and abusive content, violent and graphic content, child exploitation material, illegal and inappropriate commercial activity, dangerous organizations and individuals, false or misleading content, scams and fraud, hateful and derogatory content, sexual harassment and unwanted advances, nudity and adult content, and spam content.

User Input: LinkedIn has an active external and internal communications team that interviews users across a variety of topics and cascades member feedback to various internal teams. LinkedIn has a “feedback” button on the site for users to provide feedback. Users can also directly provide feedback about Trust issues using several in-product features like “Reporting,” “Blocking,” “I don’t know this person,” or by reaching out to LinkedIn’s customer support teams. This feedback is incorporated via various automated and manual means into LinkedIn’s proactive defenses to help improve accuracy and coverage. LinkedIn monitors metrics on these reports on a regular basis to identify and address problematic trends. These learnings are then used to inform its Trust roadmap.

LinkedIn offers members the ability to shape their Feed based on the accounts they follow. LinkedIn offers users topic-agnostic empowerment tools to control their Feed including: Disinterest, hide, and reporting content (LinkedIn alters Feed ranking based on these signals) and the ability to unfollow, block, or remove connections (which will also alter the Feed that is recommended to the member). To mitigate the impact of recommender systems, LinkedIn also allows members to sort their Feeds chronologically. LinkedIn’s Help Center also contains more information regarding how the Feed is shaped.

External Consultations: LinkedIn consults with non-governmental organizations, peer companies, and other external stakeholders to enhance policies and community standards as well as product design and enforcement. These engagements also inform the design and conduct of LinkedIn's Systemic Risk Assessment and the design of its risk mitigation measures. These engagements include the Global Internet Forum to Counter Terrorism (GIFCT), the Global Network Initiative (GNI), and the Digital Trust and Safety Partnership (DTSP). It also engages with key external stakeholders through the EU Code of Practice on Disinformation and the EU Code of Conduct on Countering Illegal Hate Speech.

LinkedIn is engaged with numerous EU bodies to address the exchange of information such as Rapid Response Group as part of the Code of Practice, the European Commission, Coimisiun na Mean, Arcom, BSI, Europol, and other pan-EU bodies. These groups provide an efficient mechanism by which LinkedIn can exchange information about issues and content that requires immediate attention.

LinkedIn works with both globally recognized independent fact checkers, human rights groups, and civil society organizations to help mitigate content- and conduct-related risks. For example, LinkedIn works with an external global news organization when its internal fact checkers are not able to confirm the veracity of user-generated content.

Document Interpretation: LinkedIn's Content Policy team, which collaborates closely with other parts of Legal as well as the Trust and Safety, Product, and Engineering teams, maintains detailed internal policies and procedures related to enforcement of the [Professional Community Policies](#). LinkedIn's internal policies are accompanied by sets of enforcement protocols that act as step-by-step guides for consistent decision-making. While the framework may vary across domains, LinkedIn leverages a common set of policy principles to define the spirit of each policy and drive ongoing research. LinkedIn documents its research efforts separately and cites them where appropriate; further, where novel policy challenges arise or precedents are set during escalations, LinkedIn leverages cross-functional issue briefs to document and refer to decision drivers.

Community Self-Regulation: Several of LinkedIn's product features allow for community-led moderation. For example, Group Admins can approve or deny member posts prior to the posts being published. Admins or Managers can also remove posts from the Group. Group Admins also have the ability to remove or block a member from the Group, which automatically removes all of their content. Members can report the post or comment, which is subject to LinkedIn's Professional Community Policies. With respect to Pages, Page admins can delete comments that are posted on their Page's posts. Admins can also add and remove other admins, depending on their admin role. Considering the professional nature of LinkedIn, self-regulation is a frequently used mechanism.

Product Enforcement

Roles & Teams: LinkedIn's Trust apparatus includes a Legal team, covering Digital Safety, Product, Platform, & Partnerships, and Privacy; Trust and Safety Operations; Engineering, Product, and Data Science, and similar teams covering unique issues related to LinkedIn Marketing Solutions and Jobs. LinkedIn has a large content moderation team with hundreds of employees and correspondingly large managed service providers that help enforce LinkedIn's Professional Community Policies and other policies. LinkedIn's Anti-Abuse team creates tools to enforce prohibition of false account activity.

Operational Infrastructure: LinkedIn operates in 36 languages (including 15 EU languages) and has 38 offices around the world. LinkedIn engages cross-functional teams in policy enforcement and maintains staff around the globe to provide linguistically and culturally relevant content moderation.

Tooling: LinkedIn has developed internal tools to allow content enforcement teams to review items that have been flagged for review by both human and automated means for potentially violating LinkedIn's Professional Community Policies. LinkedIn's flagship tool for content moderation helps reviewers look at each item (e.g., a LinkedIn post, job, ad, or message), evaluate it, and decide if it violates a LinkedIn policy. If it does, they label it appropriately and proceed to another item in the review queue.

In addition, if classifiers locate and have sufficient confidence that a particular piece of content is policy violating, LinkedIn may, depending on the severity of the violation, limit the content's visibility or remove it entirely. LinkedIn leverages content moderation technology created in-house by LinkedIn's Trust Engineering team as well as third-party enterprise tools.

Training and Awareness: LinkedIn has implemented robust training and Quality Assurance (QA) program designed to help improve reviewer performance and consistency. Regular QA audits are performed on a sample of all content items reviewed by LinkedIn's moderation teams. Results are shared weekly with more detailed summaries and reports provided monthly. Lower performers receive one-on-one coaching, common error trends are addressed during group calibration sessions, and policy refresher trainings are delivered as needed. LinkedIn has an intensive onboarding and training program for enforcement personnel, including topical training on key risk areas like elections. LinkedIn's dedicated team of trainers and quality assurance analysts are tasked with onboarding new content moderators, training content moderators on new policies and policy changes, and monitoring and improving moderator accuracy and consistency. Each reviewer goes through 3 months of training across content policies, as well as additional training as necessary to stay up-to-date or resolve areas of confusion. Content moderators also have direct access to content policy managers through regular office hours and dedicated escalation pathways. For

particularly complex decisions, content policy managers also have access to in-house lawyers who can consult regional legal experts as needed. Additionally, LinkedIn requires all employees to complete an annual mandatory compliance training that, among other things, is designed to ensure that all employees are trained in the service's approach to protecting users from harmful content or activities. LinkedIn employees are trained in the importance of data privacy and how to properly handle personal data. Finally, for GenAI features, LinkedIn maintains an internal resource that all product teams can use to find certain core legal and trust requirements, and other information about using GenAI in LinkedIn features.

Regarding external training, LinkedIn periodically runs in-product educational campaigns on LinkedIn to increase member awareness of key safety features and member control options. For example, LinkedIn provides job search safety reminders within its product to educate members on how to recognize and report suspicious jobs. This aims to increase job seeker awareness about common fraud schemes and empower LinkedIn members to (1) play a proactive role in avoiding scams and (2) contribute to LinkedIn's efforts to safeguard all members through the reporting of suspicious activities.

Finally, LinkedIn is introducing additional policy specific education and in-product interventions to educate members more comprehensively on the policies that they violate. As part of that effort, LinkedIn points members to the relevant language in LinkedIn's Professional Community Policies and Help Center articles to help members understand and contextualize in the moment why their content has been removed for running afoul of LinkedIn's policies.

Wellness and Resilience: LinkedIn Trust and Safety has a dedicated position focused on wellness and has established wellness programs in place globally to provide on demand - linguistically and culturally aligned - support to moderators who may need assistance after reviewing abusive or extreme content. LinkedIn has numerous resources available to employees that support various dimensions of wellness—from top notch medical benefits to a robust Employee Assistance Program (EAP) service, as well as comprehensive fitness and mental health programs on site and virtually. LinkedIn also has vendors in each region to support the mental health of LinkedIn content reviewers. These vendors provide group education and 1:1 wellness coaching session with a trauma-informed lens that is geared towards the most at-risk employees. LinkedIn's Trust Review OPS (TRO) team also regularly conducts "pulse checks" and surveys to understand the wellness needs of the team and shift resources as necessary to meet those needs. Separately, content moderation tooling enhancement requests are also collected and implemented on a regular basis to improve the moderator's experience, both in terms of efficacy and wellness. LinkedIn also takes into consideration the physical environment of content moderators to enhance design of working locations to promote mental wellbeing, tracks turnover and attrition as an

additional indicator of wellness and creates opportunities for team members to offboard to other roles within LinkedIn as needed.

Advanced Detection: LinkedIn uses a three-layer, multidimensional approach to moderate content within its Trust ecosystem. The first layer of protection is automatic and proactive prevention. When a member attempts to create a piece of content on LinkedIn, various calls (or signals) are sent to LinkedIn’s machine learning services. These services aim to automatically filter out certain policy-violating content within 300 milliseconds of creation, meaning the content is visible only to the author and is not shown to anyone else on the platform. As part of this process, AI tools play a key role in helping LinkedIn proactively filter out potentially harmful content. LinkedIn uses content (like certain key words or images) that has previously been identified as violating its Professional Community Policies to help inform AI models and better identify and restrict similar content from being posted in the future.

The second layer of protection is a combination of automatic and human-led detection. LinkedIn’s second layer of moderation detects content that is likely to be violative but for which the algorithm is not sufficiently confident to warrant automatic removal. This content is flagged by LinkedIn’s AI systems for further human review. The moderation management system automatically creates review items completely and accurately in the appropriate queues upon receipt of moderation requests from upstream systems. If the human review team determines that the content violates LinkedIn’s policies, it is removed from the platform. LinkedIn’s human review team is instrumental in this process and in helping train the platform’s models (as LinkedIn’s internally developed classifiers are trained on a set of confirmed policy-violating content, as reviewed by human moderators). The third layer is addressed under User Reporting. The third layer of detection is covered in the following section “User Reporting.”

LinkedIn’s GenAI features are focused on the member’s Profile and not general-purpose content creation. GenAI recommendations follow professional guidelines, enforced through various classifiers. For example, if a member prompts the AI to write a headline for their Profile including that they are an immigrant, the AI recommendation will omit that information and only include profession-relevant descriptors of the member (based on the provided prompt and the member’s Profile information).

LinkedIn applies to GenAI inputs and outputs a binary moderation process (blocking or not blocking) to minimize harmful and inappropriate content within GenAI features, including inputs and outputs containing hate, violence or sexual content. Responses are moderated in chunks of 100 tokens at a time, so -- if any portion of a response is found to be harmful or inappropriate, even if some portions of the response have already been provided -- the entire response is retracted. LinkedIn also identifies inputs relating to self-harm and redirects members who provide these inputs to self-harm resources.

LinkedIn also has in place jailbreak moderation, which is designed to detect and prevent attempts to bypass these input and output protections.

In relation to messaging, members can enable automated detection of harmful content within messages in their settings to have messages scanned for potentially harmful or harassing content.

Advertisements are automatically evaluated against LinkedIn policies at the time of their creation.

User Reporting: The third layer of LinkedIn’s content moderation approach is human-led detection. If members identify content they believe violates the Professional Community Policies, LinkedIn encourages them to report it using the in-product reporting mechanism represented by the three dots in the upper right-hand corner of the content itself or on the “More” button on the Profile headline. The reporting feature is available through, and largely identical across, LinkedIn’s website and mobile app, although reporting reasons and their visual presentation may vary slightly for certain types of content. In most instances, the reporting process is located just one click away from the content being reported and, depending on whether content is reported in the LinkedIn App or on desktop, the reporting process takes between four or five clicks to complete. Reported content is then generally enqueued for human review by LinkedIn’s Trust Review Operations team or by LinkedIn’s automated system and is removed or restricted if found to be in violation of LinkedIn’s policies. The automated system is based in part on past decisions human reviewers have made regarding whether content violates LinkedIn’s policies. When users receive notification that their report has been resolved, the notice indicates whether the report was resolved by human review or LinkedIn’s automated system. Reporting is available for all user-generated content on LinkedIn. LinkedIn makes information about user reporting easily accessible in its Help Center: [Report Inappropriate Content, Messages, or Safety Concerns](#). This page makes clear to prospective reporters that LinkedIn generally will not disclose their information with the other impacted users.

Enforcement Prioritization: LinkedIn utilizes a tool to manage prioritization of content for review. For member reports, LinkedIn’s prioritization system utilizes the predicted severity of the content, the potential for high unique member views, the report volume, number of views, and classifier scores to help ensure LinkedIn reviews the most urgent cases as quickly as possible. Other factors considered when prioritizing reviews are language, region of the originating content, and type of media. As guardrails, LinkedIn has a series of policy-based SLAs within which it aims to complete all related reviews, which are shorter for more extreme content violations (such as Terrorism and CSAM). LinkedIn also enables members to report content that portrays an intent to commit suicide / self-harm or violence against others. These reports receive priority review and are evaluated by specialists to determine how credible and imminent the threat is. When

a report is evaluated as credible and imminent, LinkedIn will attempt to notify appropriate law enforcement authorities.

LinkedIn has set, and monitors its performance against, performance targets related to the time taken to review and action flagged content, and the accuracy of content moderation decisions. LinkedIn measures its performance against its targets by tracking performance against SLAs and quality assurance scores. SLAs and quality assurance score targets are set based on a variety of factors, including minimizing the availability of illegal or policy-violating content while also making accurate content moderation decisions.

LinkedIn's Professional Community Policies apply to all members. Repeated offenses may result in feature restrictions. For certain egregious violations of LinkedIn's Professional Community Policies (e.g., child sexual abuse material, terrorism, extremely violent content, egregious sexual harassment), LinkedIn may permanently restrict an account after a single violation. Users can leverage the searchable Help Center for additional details about LinkedIn policies and how they are enforced, including how to report abusive content. LinkedIn monitors edge cases for sensitive policy escalations to document decisions and recommend policy adjustments as needed.

Appeals: LinkedIn allows users impacted by enforcement measures to appeal the relevant decision. Adequate resources are devoted to appeals review to ensure that these performance targets are regularly met. If LinkedIn reverses an enforcement decision because of an appeal, it takes steps to reverse the enforcement action. For example, content that has been removed in error is restored to the service, and restrictions that have been placed in error on accounts are likewise lifted, restoring the account to the status it held before the reversed enforcement decision. Where the error is part of a larger pattern of errors, LinkedIn uses this information to train its automated content moderation technology and/or update its enforcement guidance as appropriate.

External Reporting: In accordance with applicable legal obligations, LinkedIn reports certain CSAM-related information to NCMEC and provides member notifications in the event of significant privacy breaches. As noted in greater detail in the "Product Transparency" section below, LinkedIn also publishes information regarding its responses to government removal requests. It also voluntarily discloses certain information to relevant law enforcement authorities, as permitted by law, in the event LinkedIn becomes aware of certain imminent threats (including, for example, imminent threats of self-harm and election-related influence operations).

Third Parties: LinkedIn works with both globally recognized independent fact checkers and human rights groups to mitigate content- and conduct-related risks. For example, LinkedIn works with an external global news organization when its internal fact checkers are not able to confirm the veracity of user-generated content. LinkedIn also works with civil society groups to understand how its platform can be abused to harm members,

particularly in conflict zones. Furthermore, LinkedIn works with peer technology companies to exchange information about the tactics, techniques, and procedures (TTPs) of sophisticated abusive accounts and advanced persistent threat (APT) actors.

LinkedIn has provided a [Researcher Access Program](#) which allows independent qualified researchers access to data that is publicly accessible on LinkedIn. LinkedIn selects APIs available for approved research projects that provide access to advertising and public post data, and as appropriate, supports research requests for other datasets containing public-facing LinkedIn data.

Product Improvement

Effectiveness Testing: LinkedIn has a robust Quality Assurance (QA) program designed to help improve reviewer performance and consistency. Regular QA audits are performed on a sample of all content items reviewed by LinkedIn’s moderation teams. Results are shared weekly with more detailed summaries and reports provided monthly. Lower performers receive one-on-one coaching, common error trends are addressed during group calibration sessions, and policy refresher trainings are delivered as needed. And moderators have direct access to global content policy managers through regular office hours and dedicated escalation pathways. For particularly complex decisions, content policy managers also have access to in-house lawyers who can consult country law experts as needed.

LinkedIn measures its preventive defense services regularly to improve accuracy in the filtering process. This is done by sending some positive samples for human review to measure the precision of LinkedIn’s automated defense system. This reduces the likelihood that LinkedIn’s auto-filtering process removes content that complies with LinkedIn’s policies. LinkedIn also conducts fairness checks on Feed distribution and ranking, health checks, and long-term hold outs to make sure the algorithms are working as intended. LinkedIn looks at prevalence metrics via human review and will change the algorithm based on rising and dropping trends on the platform.

LinkedIn conducts both vendor-led and internal red-teaming (both manual and automated) efforts to test GenAI capabilities for various harms and identify mitigations for such harms. Each product team launching a GenAI feature is also responsible for undertaking quality assurance testing on response quality, accuracy, and hallucinations (where AI confidently provides a wrong or unsupported response), with goals to remediate and prevent further inaccuracies.

LinkedIn continues to monitor incidents, user reports, and social feedback to ensure that LinkedIn’s policies and principles address emerging risks and trends. A dedicated Anti-Abuse team researches emerging trends and key risks on a bi-weekly basis. One-way trends are tracked through metrics relating to user feedback. LinkedIn monitors metrics on member feedback reports regularly to identify and address problematic trends. LinkedIn also continues to monitor incidents, user reports, and social feedback to ensure that LinkedIn’s policies and principles address emerging risks and trends.

Process Alignment: A number of factors figure into LinkedIn planning and process alignment for continuous improvement. LinkedIn has both a top-down and a bottom-up approach. Trust product management sets out priorities to address risks as part of the planning process, and in addition, product managers across the organization propose solutions to emerging issues as they arise. Trust planning is a robust process. The success of LinkedIn’s large and cross-functional Trust Product team is measured on their ability to lower UMI (unique members impacted) related to content- and conduct-

related risks. In addition, LinkedIn regularly conducts manual reviews to estimate the prevalence of violative content on the site across abuse types, categorized by policy type, and uses this metric to determine the relative priority for resource allocation.

Additionally, LinkedIn continues to mature its crisis response playbook by continually monitoring crisis situations globally, expanding internal teams that work on crisis response, and maturing our processes to respond more efficiently and effectively to crisis situations. LinkedIn will continue to follow its processes related to the removal of misinformation and continually increase investments in resource allocation and process improvements where necessary to respond to the demands of the crisis.

Resource Allocation: LinkedIn rolled out a more sophisticated metric to measure the impact of risk, which can be used to allocate resources to address these systemic risks. LinkedIn has an active external and internal communications team that cascades member feedback to various internal teams within its broader Trust organization. It monitors metrics on these reports on a daily/weekly basis to identify and address problematic trends. LinkedIn also works with civil society groups to understand how its platform can be abused to harm members, particularly in conflict zones.

Additionally, in advance of each high priority election, LinkedIn reviews and determines whether an increase in resourcing is required to meet the demands of the upcoming election response. LinkedIn developed a quantitative and qualitative approach to assessing the priority of global elections in order to adequately and consistently stand up enhanced internal coverage of election risk based on relative impact. This quantitative and qualitative assessment is used to determine the length and severity of the risk mitigation measures that need to be implemented for each high priority election, including the language resourcing requirements. Depending on the election, risk mitigation measures at LinkedIn are in place and functioning one to six months before an electoral period and can continue for one month after the elections.

To ensure adequate resourcing at times of crisis, in addition to ad hoc tiger teams, LinkedIn has developed incident response plans and escalation paths to address potential crisis events. LinkedIn employs a “follow the sun” model for crisis events and has created a Follow the Sun Escalation directory and provides training on how to leverage the directory.

External Collaboration: LinkedIn consults with non-governmental organizations, peer companies, and other external stakeholders to enhance policies and community standards as well as product design and enforcement, including the Global Internet Forum to Counter Terrorism (GIFCT), the Global Network Initiative (GNI), and the Digital Trust and Safety Partnership (DTSP).

As a member of GIFCT, LinkedIn has access to GIFCT’s Incident Response processes, including hashes related to an event activated as Content Incidents or Content Incident

Protocols. This allows LinkedIn to quickly become aware of, assess, and address potential content circulating online resulting from a terrorist or violent extremist event. For further information, see GIFCT's annual Transparency Report, which includes information on the hash-sharing database. LinkedIn also works with GIFCT and fellow members of that organization to solicit input on their content policies.

Remedy Mechanisms: Members can submit general complaints to LinkedIn by creating a support ticket in LinkedIn's Help Center. Members can submit Privacy complaints through LinkedIn's Privacy Hub. Where LinkedIn removes a member's content or suspends their account, LinkedIn generally provides the member notice of the action LinkedIn has taken and provides them with an entry point through which to appeal LinkedIn's decision. LinkedIn reviews those appeals and notifies the member of the appeal decision.

Training & Awareness: Educational campaigns curated by LinkedIn communications team through LinkedIn's transparency center, LinkedIn's blog and LinkedIn's Company Page posts help to bring awareness to the reporting mechanism. In product notices, such as inline warnings, are also used to encourage reporting of content that may violate LinkedIn's policies, as detected by LinkedIn's automated defenses.

Product Transparency

Transparency Reporting: A cross-functional team of stakeholders from across LinkedIn's broader Trust organization (including members of its Legal, Trust Review Ops and Data Science and Communications teams) work closely to produce a semiannual global Transparency Report. LinkedIn's global [Transparency Report](#) is issued twice per year and is published in LinkedIn's Transparency Center ([linkedin.com/transparency](#)). Additional country-specific Transparency Reports are also published depending on local requirements. LinkedIn's Transparency Reports currently report on the following metrics: fake accounts (including the number removed and percentages stopped at different stages); spam and scams (including percentages stopped by automated defenses and number removed proactively/after member reports); content removed under its Professional Community Policies (including harassment or abusive, misinformation, hateful or derogatory, violent or graphic, adult, and child exploitation); and copyright removals (including number of requests, total infringements reported, reported infringements removed/rejected). LinkedIn's Transparency Reports also include certain data on the number and types of government requests received and actioned by LinkedIn (including requests for user data and for content removal).

LinkedIn also publishes a [global Ads Library](#) where users can access information about ads that have run on LinkedIn. LinkedIn's 'Why am I seeing this ad' feature also includes additional information about the type of ads members see and other options to manage the advertising experience.

Notice to Users: When a member's content is removed or either their account or access to certain LinkedIn features is restricted because their content does not comply with LinkedIn's policies, LinkedIn generally provides the member notice and the opportunity to appeal the decision except in certain limited circumstances like, for example, where the content at issue is deceptive high volume commercial content (e.g., fraudulent) or LinkedIn is prohibited from providing notice. Notices are typically sent by email and contain a link to a notice page. LinkedIn reviews those appeals and notifies the member of its appeal decision. At the time a content creator (including ad publishers) is notified their content was moderated or restrictions were applied to their Profile, LinkedIn publishes the statement of reasons, which do not include personal data, to the European Commission's (EC's) DSA database.

Where content does not violate LinkedIn's policies, but LinkedIn removes or geo-blocks the content in response to valid legal process demanding its removal, LinkedIn generally provides prompt notice to the member(s) whose content LinkedIn has actioned, unless legally prohibited from doing so. This notice allows the impacted member to challenge the legal order served on LinkedIn.

LinkedIn's terms of service are updated periodically, as needed. Material updates to the User Agreement and Privacy Policy are communicated directly to members – by means of emails or in-product notifications in advance of the update, together with a preview of the update – and LinkedIn provides persistent links that allow users to access the terms.

Additionally, when a user reports content through one of its dedicated reporting channels, LinkedIn acknowledges receipt of the report by email. If the reporter uses LinkedIn's in-product reporting flow to report content as policy violating, LinkedIn's email confirming receipt will contain a link to a notice page and instructions advising that the reporter can use the link to find information about the current status of their report (including LinkedIn's decision whether to action the reported content). LinkedIn updates the notice page in near real-time. LinkedIn members also can access this notice page at any time by visiting their "Report History" in their Notification settings.

Complaint Intakes: Members can submit general complaints to LinkedIn by creating a support ticket in LinkedIn's Help Center. Members can submit Privacy complaints through LinkedIn's Privacy Hub.

Researcher & Academic Support: LinkedIn has long explored various mechanisms to support research on the issues related to LinkedIn's platform, including emerging skills and skill gaps (e.g., enabling academic research through the economic graph challenge, partnering with the World Bank). LinkedIn and Microsoft are collaborating with OpenMined on the "Initiative on Algorithms" to explore how to enable further research collaboration by exploring research infrastructure that protects user privacy and proprietary information, while achieving reproducibility, affordability, and scale for independent researchers. In furtherance with its commitments under the Code of Practice on Disinformation, LinkedIn is working with Trustlab to enable reporting on misinformation. More recently, LinkedIn has been supporting new research mechanisms related specifically to its ads and Feed, including by making a tool available that supports research into ad content shown on LinkedIn. It is also launching new mechanisms that will enable qualifying researchers to access public data (such as Feed) and ads.

Furthermore, LinkedIn supports researchers through the [\[beta\] Research Access Program](#). Researchers interested in using LinkedIn's publicly available data may fill out an [application form](#). Researcher applicants who satisfy the criteria set forth in the DSA are notified and asked to enter into terms governing their access to the public data (and/or use of APIs, as applicable) for the purpose of the approved research, including terms related to the security and legal use of personal data.

In-Product Indicators: LinkedIn uses a variety of in-product indicators to enhance member awareness and empowerment, including "Why am I seeing this ad?" as well as "Kindness Reminders" (currently deployed in English) to encourage members to think

twice before posting comments that may be harmful or violate their Professional Community Policies.

LinkedIn offers settings that allow members to only receive communications from people they know. It also offers an optional advanced safety feature in which members can allow LinkedIn's automated machine learning models to detect sexual harassment and certain other abusive content in messages received by the member, giving the recipient the ability to view and report the message to LinkedIn for review if desired. Where the setting is enabled and LinkedIn's automated systems detect potentially harmful content, depending on whether the sender and recipient have exchanged prior communications, the message will be directed to the recipient's spam folder or be hidden from the recipient by a warning. The warning can be dismissed, giving the recipient the ability to view and report the message to LinkedIn for review if desired. In certain circumstances LinkedIn also may surface a warning and fraud prevention resources to members if its technology detects an attempt to move private conversations away from LinkedIn. This warning is designed to encourage members to consider the likely authenticity of the people with whom they are engaging.

In LinkedIn's commitment to uphold Responsible AI Principles and legal frameworks, LinkedIn has implemented a set of guidelines to ensure transparency and clear labeling across all AI-driven tools, products, or features. This approach involves consistent iconography and language throughout the member experience, context-appropriate labeling to distinguish AI from non-AI-generated content, adherence to legal and trust guidelines to avoid deceptive practices, and the inclusion of feedback mechanisms for GenAI experiences. Content containing the Coalition for Content Provenance and Authenticity's industry-leading "Content Credentials" technology (C2PA) will be automatically labelled on LinkedIn. Users will see the "Cr" label. By clicking the label, users will be able to trace the origin of the AI-created media, including the source and history of the content, and whether it was created or edited by AI.

LinkedIn offers members the choice to engage with AI-driven features, accompanied by prominent AI disclosures regarding potential inaccuracies. LinkedIn encourages user review, editing, and approval of AI-suggested texts, and provides direct links for members to Help Center articles. This approach underscores LinkedIn's dedication to fostering a responsible, trustworthy digital environment, leading by example in the Responsible AI space and ensuring the platform remains a safe, transparent, and empowering space for professional growth and connection.

7.2 Monitoring the Effectiveness of Mitigations

LinkedIn has robust processes in place to monitor the effectiveness of mitigations implemented to address systemic risks on the service.

LinkedIn's Trust team maintains and regularly monitors a range of dashboards to determine the health of the platform, including dashboards tracking that the following perform as intended:

- classifiers designed to auto-action or enqueue content for human review;
- LinkedIn's hide, mute, unfollow, block, and disinterest functionalities; and
- LinkedIn's content moderation decision quality assurance audit process.

Designated teams within LinkedIn's broader Trust organization monitor these dashboards regularly to measure their performance against certain target or acceptable thresholds. Additionally, a number of dashboards are configured to send real time alerts to LinkedIn's Trust team whenever an issue with a functionality is detected. Where an alert is triggered or a dashboard's metrics deviate from acceptable levels, the Trust team investigates and remediates any issues or deviations. Trust leadership also receives regular updates on the monitoring and performance of these dashboards, providing leadership regular visibility into the ongoing effectiveness of these mitigations.

AI plays a key role in helping LinkedIn proactively filter out policy-violating content. LinkedIn uses such content to help inform AI models and better identify and restrict similar content from being posted in the future. LinkedIn's data science team monitors the content filtering process to measure how many content violations are successfully prevented and how much remains on the platform. Content that is proactively removed at creation is tracked through a data pipeline, and LinkedIn measures its preventive defense services regularly to improve accuracy in the filtering process. This is done by sending a set of positive samples for human review to measure the precision of LinkedIn's automated defense system.

The Trust Review Operations team also regularly reviews content moderation metrics to monitor the efficacy of its mitigations and to identify where additional mitigations may be needed to account for new abuse vectors. LinkedIn regularly conducts manual reviews to estimate the prevalence of violative content, or Policy Violative Views (PVV) on the site across abuse types, categorized by policy type, and uses this metric to determine the relative priority for resource allocation and to assess the overall effectiveness of its mitigation measures. The overall average percentage of user views of violative content - as compared to views of non-violative content - on the LinkedIn platform during the assessment period was 1.23%. The average percentage of user views of violative content in Feed was 1.4%; views of violative Jobs was 0.8%; views of violative Ads was .2%; and exposure to content from violative accounts, such as connection requests, was 4.9%.

LinkedIn continually looks for opportunities to further enhance its efforts to assess the effectiveness of its mitigation measures and, as described in [Planned Enhanced Mitigation Efforts](#), is investing in further refining the alignment of PVV tracking to the systemic risk areas and more nuanced geographic regions.

LinkedIn's Trust and Safety processes are periodically audited by the Company's Internal Audit team, which conducts thorough and objective evaluations of key risks, processes, and mitigates, helping LinkedIn identify areas for improvement and maintain high standards of compliance and controls. LinkedIn also follows Management Action Plans mandated by the Internal Audit team.

7.3 Progress Made on Enhanced Mitigations

In last year's Systemic Risk Assessment report, LinkedIn outlined areas of focus for enhancement of safety mitigations during the 2024 – 2025 reporting period. LinkedIn has made enhancements in each identified area as described below:

- Identify opportunities to more directly map classifiers and internal risk taxonomy to systemic risk areas in order to enhance the monitoring of mitigations.
 - LinkedIn selected Policy Violating Views (PVV) metric as a streamlined, formalized, and easy to audit way to evidence continuous monitoring of systemic risk on the platform.
- Invest in further expanding upon the improvements made its models built to detect fake accounts to achieve greater fidelity on detection of additional inauthentic activity.
 - This is an area of continuous investment for LinkedIn, and is evidenced by its product and engineering roadmaps for Fiscal Year (FY)'25 and FY'26.
- Tuning and moderating classifiers, multilingual models, and enhanced automation capabilities, particularly for risks related to Human Dignity and Consumer Protection and Fraud.
 - This is an area of continuous investment for LinkedIn, and is evidenced by its product and engineering roadmaps for FY'25 and FY'26.
- Conduct post-mortem reviews following high priority elections to identify opportunities to tighten and refine LinkedIn's election response capabilities.
 - LinkedIn has put in place post-mortem reviews in alignment with the risks associated with the underlying election.
- Enhance protective measures to reduce the impacts of verification abuse, as appropriate, based on additional review criteria
 - LinkedIn put in place a range of anti-abuse measures to better prevent inauthentic actors from obtaining or misusing LinkedIn verifications. It also carried out member education work during the assessment cycle to help members better understand the benefits of authenticity signals like verifications.

- Invest in efforts to identify fraudulent accounts, malware, and other Tactics, Techniques, and Procedures that may be used for fraud and/or degrading consumer protection
 - This is an area of continuous investment for LinkedIn, and is evidenced by its product and engineering roadmaps for FY'25 and FY'26.
- Explore opportunities to further enhance the pre-launch product review and launch process
 - LinkedIn established and began using a more standardized Critical Impact Risk Assessment (CIRA) screening and secondary review process.
- Continue investing in and being transparent in its practices related to researcher access to data to researchers, members, and the public.
 - LinkedIn is currently working with at least one approved researcher to identify relevant data sets that can be made accessible to them and to establish appropriate terms and other safeguards to protect such data.

7.4 Planned Enhanced Mitigation Efforts

In accordance with DSA Article 35, LinkedIn has included in [Appendix B](#) a list of focus areas for the coming year to further enhance the reasonable, effective, and proportionate mitigation of systemic risks stemming from the use, misuse, or functioning of the LinkedIn platform. The mitigation focus areas include, for example:

- Continued tuning and moderating of classifiers, multilingual models, and enhanced automation capabilities, for risks related to Consumer Protection and Fraud and Human Dignity.
- Evaluation of mitigations in place for online protection of minors and potential enhancements thereto as appropriate following LinkedIn's review of the suitability of the mitigations set forth in the newly published Article 28 Guidelines.
- Continued investment in tracking prevalence metrics to measure the effectiveness of LinkedIn's mitigation measures.
- Continued enhancements to and refinement of detection and removal of non-consensual intimate imagery.

8. Conclusion

The objective of the Systemic Risk Assessment is to assess areas of greater potential risk stemming holistically and objectively from the function, use, or misuse of the LinkedIn platform. This exercise enables LinkedIn, as a Very Large Online Platform, to better plan for and prioritize enhancements related to people, policies, processes, and technologies that will further reduce risk of negative impact on the fundamental rights of users. While overall systemic risk related to LinkedIn's products and services are assessed as low, LinkedIn has leveraged this assessment to identify opportunities to continue refining, enhancing, and expanding its efforts to keep its platform and services safe, trusted, and professional.

9. Appendices

Appendix A: Methodology Expanded

Appendix B: List of Article 35 mitigations

Appendix A: Methodology Expanded

Risk Areas

LinkedIn defined 12 key Risk Areas to facilitate an assessment of systemic risks stemming from the use or misuse of the LinkedIn platform or services as required by DSA Article 34. Risk Scenarios are used to further describe the risks that may fall under each risk area. Not all Risk Scenarios are applicable to each LinkedIn feature.

Figure 35: Key Risk Areas

Risk Area	Risk Definition	Risk Scenarios
Civic Discourse and Electoral Processes	Risk that content or activities negatively impacting civic discourse or electoral processes occur on the platform.	- Coordinated manipulation, disinformation, inauthentic activity, and foreign interference. - Election misinformation. - Political polarization, political bias, and the creation of echo chambers.
Consumer Protection and Fraud	Risk that activities that are fraudulent or undermine consumer protection occur on the platform.	- Scamming, defrauding, spamming, fraudulent business or profile. - False representation, abuse of position, falsifying information, including the creation of fake accounts, use of bots, or other deceptive commercial practices.
Discrimination and Hate	Risk that discriminatory content or activities that target or exclude individuals, groups, or businesses based on protected grounds or promote bias or discriminatory practices occur on the platform.	- Bias and discriminatory practices. - Hate speech and discriminatory content including racially or religiously aggravated public order offences and harassment.
Freedom of Expression and Information	Risk that freedom of expression, information, the press, or pluralism of the media are negatively impacted by the design or functioning of a platform and its algorithmic systems, policies, or enforcement.	- Over moderation and methods for silencing expression. - Recommendations that suppress pluralism.
Human Dignity	Risk that content or activities degrading human dignity occur on the platform.	- Human trafficking (other than sex trafficking) or coercive conduct depriving others of their human rights for personal gain. - Sexual exploitation, non-consensual prostitution or pornography, or sex trafficking for personal gain. - Graphic violence, human gore, vulgarity, profanity, or other content that demeans, shocks, or offends.

Risk Area	Risk Definition	Risk Scenarios
		- Sexual or pornographic content, and consensual, legal prostitution. - Extreme pornography.
“Illegal” Content and Activities	Risk that illegal content or activities are shared or perpetrated on the platform.	- Intellectual property infringement, defamation. - Promotion, concealment, acquisition, supply or sale of drugs or psychoactive substances and services relevant to the same. - Facilitating unlawful immigration and illegal entry. - Causing suffering of animals. - Creation, dissemination, or promotion of Child Sexual Exploitation and Abuse Imagery (CSEAI) including image-based CSAM and CSAM URLs. - Promotion, concealment, acquisition, use or sale of criminal property and services relevant to the same. - Promotion, concealment, acquisition, or sale of firearms, and other weapons including ammunition, crossbow, knives and services relevant to the same.
Mental and Physical Well-being	Risk that activities potentially causing harm to mental or physical wellbeing occur on the platform.	- Gender based violence. - Threats of violence or to kill, harassment, stalking, public order offenses, cyber-bullying and fear or provocation of violence. - Encouraging or assisting suicide or attempted suicide or serious self-harm. - Behavioral addictions impacting mental health. - Controlling or coercive behavior.
Private and Family Life	Risk that content or activities negatively impacting the right to private and family life occur on the platform.	- Malicious sharing or exploitation of sensitive or private data and doxing. - Disclosing or threatening to disclose nonconsensual private or intimate images (NCII). - Privacy intrusions.
Protection of Personal Data	Risk that activities that negatively impact protection of personal data occur on the platform.	- Collection or processing of information without consent or other misuse of data, including data mining or targeted advertisements without consent. - Data breaches or other unauthorized disclosure or insufficient protection of data.

Risk Area	Risk Definition	Risk Scenarios
Public Health	Risk that content or activities negatively impacting public health occur on the platform.	• Promotion of legal but harmful substances or practices. • False or misleading health related information or misinformation.
Public Security	Risk that content or activities degrading public security occur on the platform.	• Terrorist recruitment, funding, or training, terrorist imagery or content, or support for terrorist acts. • Mis- or disinformation undermining public security institutions or related to crisis events.
Rights of and Protection of Minors	Risk that content or activities negatively impacting the rights, privacy, safety, and security of minors occur on the platform.	• Child sexual exploitation and abuse and grooming. • Collection of children's data without parental consent or targeting children with personalized advertisements. • Exposure of children to harmful content or activities, including gambling, exploitation, crime, violence, terrorism, or violent extremism. • Discrimination, cyberbullying, and harassment of minors. • Promotion of self-harm or self-injury for minors.

Assessment Process

LinkedIn followed the DTSP Safe Assessment framework process to guide the conduct of the risk assessment.

DTSP Safe Assessment step descriptions

Figure 36: DTSP Safe Assessment Process

Step	Description	Objective
Discover relevant information	Engage key product stakeholders and perform initial information discovery on the company's practices across the 5 DTSP commitments and identify the practices to be evaluated for their use in mitigating content and conduct risks.	Establish baseline understanding of the operational landscape and identify the specific DTSP practices used to mitigate content- and conduct-related risks.
Identify and prioritize relevant risk considerations	Using the artifacts and information collected during the "Discover" stage - identify, document, and prioritize risks about the ways that content- and conduct-related risks are identified and mitigated.	Prioritize risks about the ways that content- and conduct-related risks are identified and mitigated to inform focus areas for the assessment.
Assess practices and risk mitigation	For the relevant risks about the ways that content- and conduct-related risks are identified and mitigated at the company and focus areas identified in the previous step, analyze the practices employed to control for, or protect against, Trust and Safety risks.	Understand current state processes, practices, and tools in relation to a common maturity scale.
Test control strength and effectiveness [Level 2 and Level 3 only]	Perform a control strength evaluation, including control design and effectiveness testing.	Understand, at a granular level, the operational effectiveness of risk mitigation processes, procedures, and tools.
Report results and findings	Compile all analysis results and report out on findings, observations, and future opportunities for improvement on the ways that content- and conduct-related risks are identified and mitigated at the company moving forward.	Share key observations and findings with partners to facilitate collaborative development of industry standards and perspective.

Risk Assessment Inputs

LinkedIn gathered and reviewed the following inputs to inform and substantiate ratings assigned in the systemic risk assessment.

- **DTSP Questionnaire Responses:** LinkedIn utilized the modified April 2023 Digital Trust and Safety Partnership (DTSP) questionnaire to solicit responses to the questionnaire from internal stakeholders across the Trust and Safety, Legal, Privacy, Product, and Engineering teams.
- **Product and Feature Changes:** LinkedIn collected material changes to in-scope features through the Critical Impact Risk Assessment (CIRA) process during the assessment period.
- **Internal Consultations:** LinkedIn conducted in-depth workshop sessions with internal stakeholder groups to solicit more detailed information on specific topics or issues related to risk prevalence as well as best practices and mitigating factors with greater potential to impact LinkedIn's risk profile.
- **External Consultations:** LinkedIn regularly engages, both individually and with parent company Microsoft, with external stakeholders, including civil society organizations, to receive feedback on platform Trust and Safety as well as to discuss best practices for addressing risk. LinkedIn conducted additional civil society engagement specific to this systemic risk assessment and the feedback from these engagements informed the assessment.
- **Policies and Publications:** LinkedIn reviewed relevant external policies, public practices, and other authoritative sources including blogs to identify additional policies and initiatives most relevant to the risk assessment.
- **Authoritative Sources:** LinkedIn reviewed publicly available sources considered reliable due to their expertise and reputation, such as regulatory sources, reputable public opinion polling, and think tank research, to inform the objective assessment of severity.
- **Transparency Report Metrics:** LinkedIn considered metrics reported through Transparency Reporting to inform the assessment of both probability and effectiveness of mitigations.
- **Open-Source Data on Public Discourse:** LinkedIn reviewed a collection of social and digital media articles and conversations concerning LinkedIn and each systemic risk area to identify trends in areas of public discourse and/or concern.
- **Internal Content Moderation Metrics:** LinkedIn considered internal metrics related to content moderation to inform both assessment of prevalence as well as effectiveness of mitigations.
- **DSA Risk and Control Matrix:** LinkedIn considered the DSA risk and control matrix as well as the results from the DSA audit to inform the existence and effectiveness of mitigations and Residual Risk scores.

Inherent Risk

Inherent risk is determined by two factors - probability and severity. Calculating inherent risk involves assessing both the likelihood of an event occurring and the potential impact of that event. The greater the probability and severity of an event, the higher the inherent risk associated with it.

Probability: LinkedIn conducted a **data-driven probability assessment** to evaluate the likelihood of certain events occurring on the platform absent mitigations by analyzing relevant data, including public incident data, transparency report metrics, and internal metrics.

- Probability is the likelihood of the probable worst case of the risk occurring if management action and controls / mitigations fail or do not exist.
- Probability depends on factors such as volume of content flagged or removed for identified risk events.
- Each score assigned to probability in calculating Inherent Risk, ranging from 1 to 5, corresponds to a specific range of likelihood of the event occurring, with higher scores indicating a higher likelihood of the event.
- Items that substantiate probability include External consultations, Internal content moderation metrics, Transparency reporting metrics, and Open-source data on public discourse related to LinkedIn and systemic risk areas.

Figure 37: Probability Rating Scale

Description	Score	Rating
The risk event or circumstance is highly likely to occur without considering mitigations in place.	5	Expected
The risk event or circumstance is highly likely to occur without considering mitigations in place.	4	Highly Likely
The risk event or circumstance is likely to occur without considering mitigations in place.	3	Likely
The risk event or circumstance occurring is possible but not likely without considering mitigations in place.	2	Not Likely
The risk event or circumstance is only remotely probable without considering mitigations in place.	1	Remote

Severity: LinkedIn conducted an objective systems-based analysis to determine severity, considering the complexity, scale, and gravity of impact to assign an overall severity rating.

- Severity is the probable worst case of a risk's impact on the world or society if management action and controls / mitigations fail or do not exist.
- Severity is calculated once for each risk; and takes consideration of complexity (impact on economic, security, political, societal, wellbeing, and environmental systems), scale (global, regional, country, locality, individual), and gravity (remediability)
- Items that substantiate severity include External consultations, Internal consultations, and Authoritative sources.

Figure 38: Severity Rating Scale

Description	Weight		Score	Rating
	Min	Max		
Impact that could cause critical, irremediable harm, damage, or loss.	145	270	5	Critical
Impact that could cause significant irremediable harm, damage, or loss.	95	145	4	High
Impact that could cause some harm or disruption but is generally manageable or remediable.	75	95	3	Moderate
Impact that could cause limited harm or disruption.	71	75	2	Low
Impact that has little or no consequence.	70	71	1	Minimal

Figure 39: Inherent Risk Rating Scale

Description	Score	Rating
A risk that is relatively certain to occur, would have a severe impact if it occurred, and requires immediate action to manage or mitigate.	17 <=	Critical
A risk that is highly likely to occur, would have a significant impact if it occurred, and requires urgent action to manage or mitigate.	11-17	High
A risk that is likely occur, would have a noticeable impact if it occurred, and requires some action to manage or mitigate.	6 - 11	Moderate
A risk that is not likely to occur, would have a limited impact if it occurred, and may require minimal action to manage or mitigate.	2 - 6	Low
A risk that is only remotely probable, would have a minimal impact if it occurred, and may not require any action to manage or mitigate.	0 - 2	Minimal

Residual Risk

Residual Risk: By multiplying the Inherent Risk score and the Maturity of Mitigations score, a Residual Risk score is calculated that reflects the amount of risk remaining in a given situation. This provides a final assessment of the potential risk once mitigations have been applied.

Maturity of Mitigations: LinkedIn leverages the DTSP Safe Assessment Framework – aligning controls and mitigations to the 35 Best Practices – and Maturity Scale for mitigations, while evaluating if they are reasonable, proportionate, and effective.

- LinkedIn first assesses whether there are mitigations in place to address each type of risk on each feature where that risk manifests.
- If there are no gaps in risk coverage, LinkedIn then evaluates the volume or depth of unique risk mitigation measures for each risk area.
- Finally, LinkedIn considers the breadth of LinkedIn’s mitigation measures assigned to each risk area, assigning a higher score where all DTSP Best Practice categories are covered.
- Items that substantiate mitigation include DTSP questionnaire responses, Internal consultations, External consultations, Internal content moderation metrics, Mitigation summaries, and Policies and publications.

Figure 40: DTSP Inventory of 35 Best Practices

DTSP Inventory of 35 Best Practices				
Product Development	Product Governance	Product Enforcement	Product Improvement	Product Transparency
PD1: Abuse Pattern Analysis	PG1: Policies & Standards	PE1.1: Roles & Teams	PI1: Effectiveness Testing	PT1: Transparency Reports
PD2: Trust & Safety Consultation	PG2: User Focused Product Management	PE1.2: Operational Infrastructure	PI2: Process Alignment	PT2: Notice to Users
PD3: Accountability	PG3: Community Guidelines/Rules	PE1.3: Tooling	PI3: Resource Allocation	PT3: Complaint Intakes
PD4: Feature Evaluation	PG4: User Input	PE2: Training & Awareness	PI4: External Collaboration	PT4: Researcher & Academic Support
PD5: Risk Assessment	PG5: External Consultation	PE3: Wellness & Resilience	PI5: Remedy Mechanisms	PT5: In-Product Indicators
PD6: Pre-Launch Feedback	PG6: Document Interpretation	PE4: Advanced Detection		
PD7: Post-Launch Evaluation	PG7: Community Self Regulation	PE5: User Reporting		
PD8: User Feedback		PE6.1: Enforcement Prioritization		
PD9: User Controls		PE6.2: Appeals		
		PE6.3: External Reporting		
		PE7: Flagging Processes		
		PE8: Third Parties		
		PE9: Industry Partners		

https://dtspartnership.org/wp-content/uploads/2022/07/DTSP_Report_Safe_Assessments.pdf

Figure 41: Maturity of Mitigations Rating Scale using DTSP Maturity Descriptions

Description	Risk Coverage	Framework Coverage	Final Score	Rating
A rating of Ad Hoc is assigned when execution of best practices is incomplete, informal, or inconsistent.	0	N/A	10%	Ad Hoc
A rating of Repeatable is assigned when execution of best practices occurs without standardized processes. Organizations aim to document more formalized practices.	14	N/A	30%	Repeatable
A rating of Defined is assigned when execution of best practices occurs with defined and documented processes. Processes are more proactive than reactive and are implemented across the organization. A risk can only receive a Defined coverage score if there are mitigations in place on each product to address all risks relevant to that product.	21	N/A	50%	Defined
A rating of Managed is assigned when execution of best practices is defined, documented, and managed through regular reviews. Organizations use feedback to continuously mitigate process deficiencies. A risk can only receive an optimized score if, in addition to having full risk coverage, the mitigations in place cover 65% of framework categories.	28	50%	65%	Managed
A rating of Optimized is assigned when execution of best practices promotes Trust & Safety in every aspect. Processes are continuously improved with innovative ideas and technologies. A risk can only receive an optimized score if, in addition to having full risk coverage, the mitigations in place cover 100% of framework categories.	N/A	100%	75%	Optimized


https://dtspartnership.org/wp-content/uploads/2022/07/DTSP_Report_Safe_Assessments.pdf

Figure 42: Residual Risk Rating Scale

Description	Score	Rating
A risk that is relatively certain to occur would have a severe impact if it occurred and requires immediate action to manage or mitigate.	17 - 32	Critical
A risk that is highly likely to occur, would have a significant impact if it occurred, and requires urgent action to manage or mitigate.	11 - 17	High
A risk that is likely occur, would have a noticeable impact if it occurred, and requires some action to manage or mitigate.	6 - 11	Moderate
A risk that is not likely to occur, would have a limited impact if it occurred, and may require minimal action to manage or mitigate.	2 - 6	Low
A risk that is only remotely probable, would have a minimal impact if it occurred, and may not require any action to manage or mitigate.	0 - 2	Minimal

Sample Scoring Calculation

Figure 43: Sample Scoring Calculation



Inherent Risk

- If the probability of occurrence of a risk is Highly Likely -> probability score of 4
- If the severity of a risk is Moderate -> severity score of 3
- Based on the probability score of 4 and severity score of 3 -> Inherent Risk score of 12 (= 4 x 3) -> inherent rating of High

Residual Risk

- If the inherent risk rating score is 12
- If the mitigating maturity rating is "Managed" -> mitigation score of 65%
- Based on the inherent rating score of 12 and mitigation score of 65% -> Residual Risk score of 4.2 (= 12 x (1 – 0.65)) -> Residual Risk rating of Low

Appendix B: List of Article 35 Mitigations

The following areas have been identified for heightened focus for the coming year, in addition to the ongoing risk mitigation monitoring and refinement processes. LinkedIn is committed to addressing each of the topics identified below to (1) further define the associated risks, (2) explore enhanced mitigation options, (3) vet mitigation options as appropriate with internal and external subject matter experts, (4) develop action plans as needed to address focus areas, and (5) track implementation of the developed action plans.

Figure 44: Article 35 Mitigations

Mitigation Title	Mitigation Description
Refined Models to Detect Additional Fake Behaviors	LinkedIn will invest in further expanding upon the improvements made its models built to detect fake accounts to achieve greater fidelity on detection of additional inauthentic activity.
Continuous Tuning of Classifiers	LinkedIn plans to continue tuning and moderating its classifiers, multilingual models, and enhanced automation capabilities, particularly for risks related to Human Dignity and Consumer Protection and Fraud.
Fraudulent Accounts	LinkedIn plans to invest in efforts to identify fraudulent accounts, malware, and other Tactics, Techniques, and Procedures that may be used for fraud and/or degrading consumer protection.
Rights and Protection of Minors	LinkedIn plans to evaluate whether the mitigation measures proposed in the European Commission's recently published Article 28 Guidelines are appropriate and proportional to ensure the privacy, safety and security of minors on LinkedIn. Based on its findings from that evaluation, LinkedIn may refine or supplement its existing compliance measures for online protection of minors in alignment with Article 28.
Continued Investments in Prevalence Metrics	LinkedIn plans to continue investing in tracking prevalence metrics to measure the effectiveness of its mitigation measures.
Non-consensual Intimate Imagery	LinkedIn plans to focus on continued enhancements to and refinement of detection and removal of non-consensual intimate imagery.